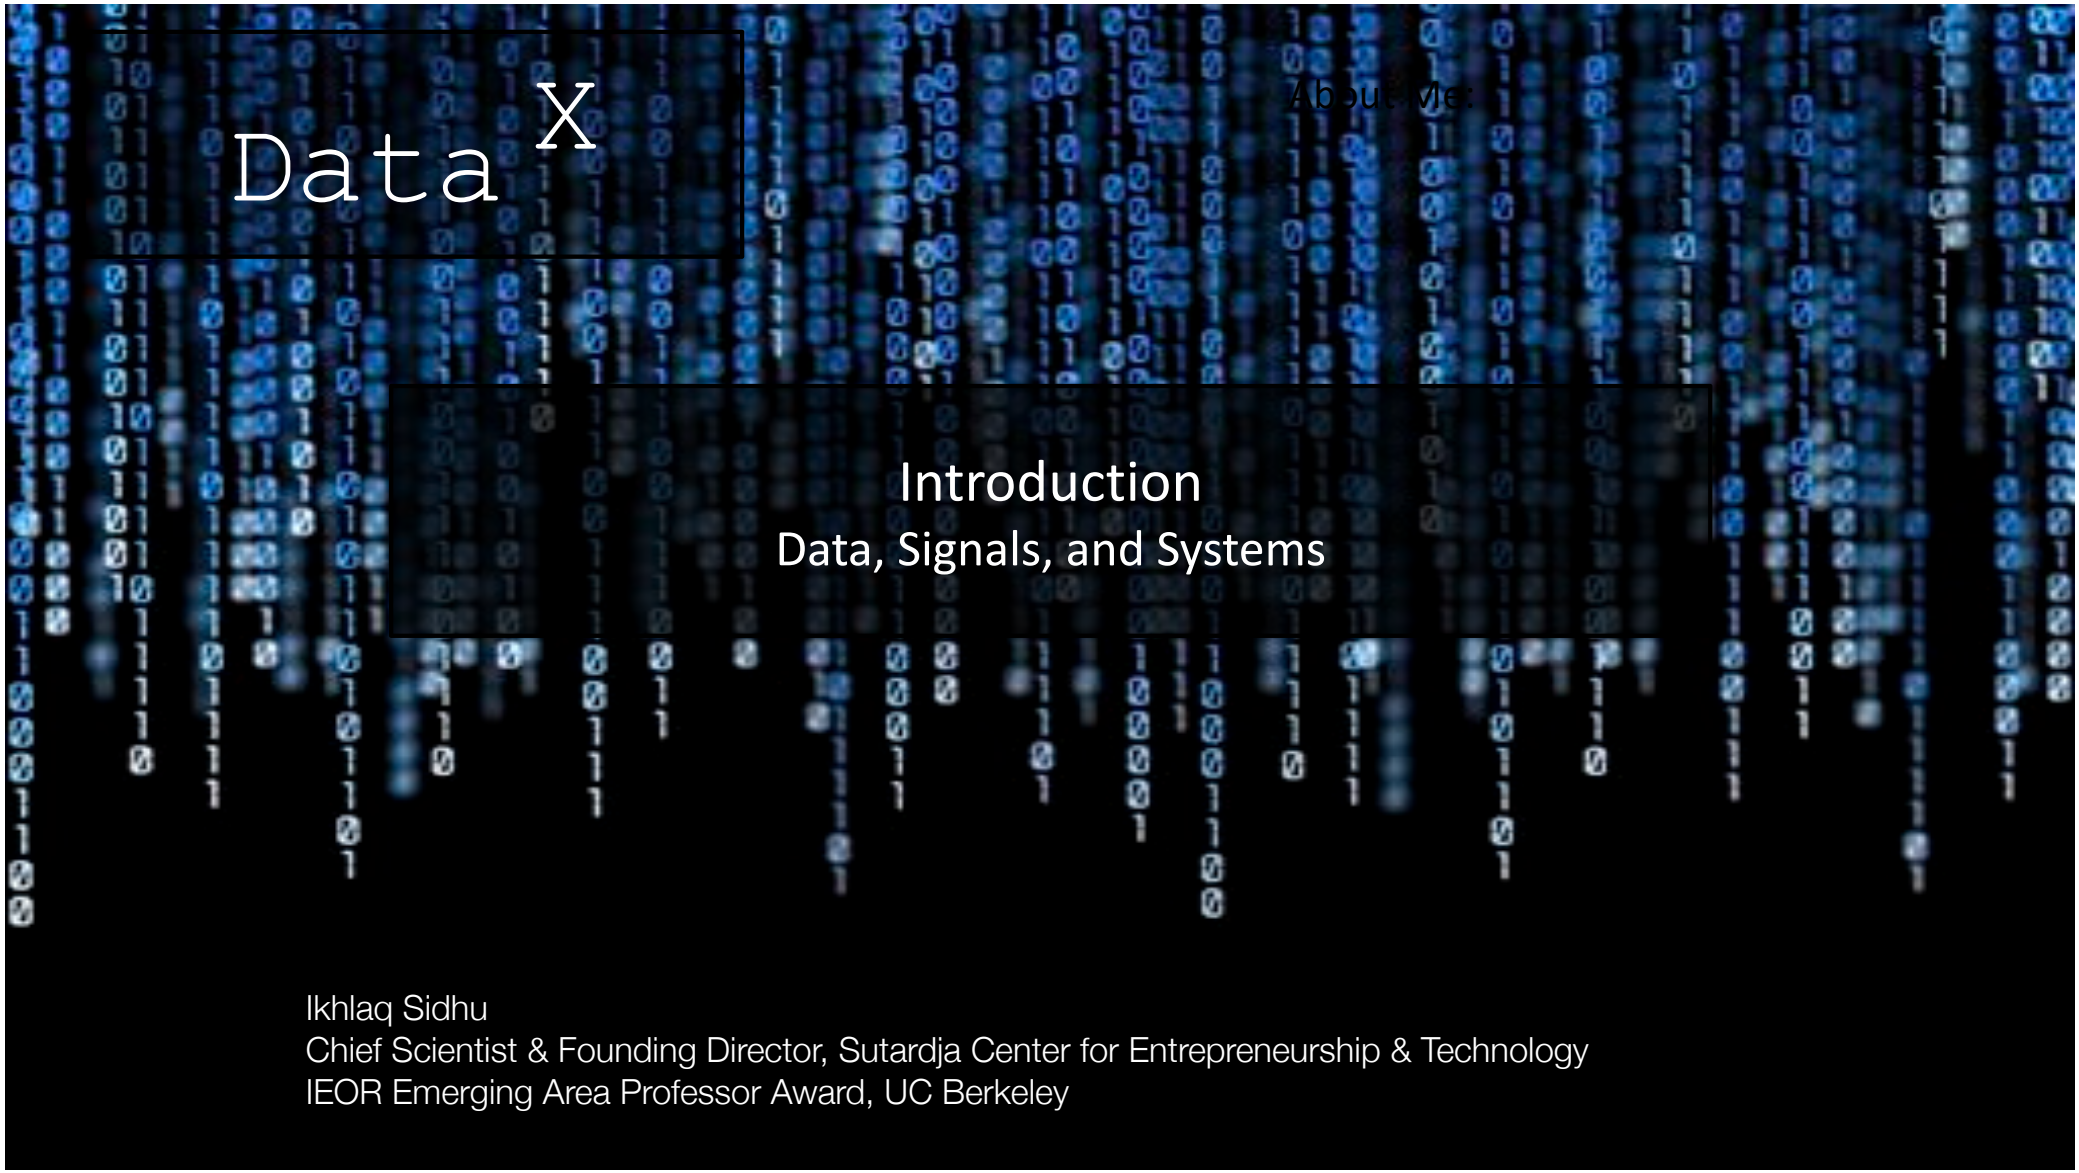
Data X
About Me:
Introduction
Data, Signals, and Systems
Ikhlaq Sidhu
Chief Scientist & Founding Director, Sutardja Center for Entrepreneurship & Technology
IEOR Emerging Area Professor Award, UC Berkeley

Welcome to Applied Data Science with Venture Applications

Sutardja Center for Entrepreneurship & Technology
Industrial Engineering End Operations Research Department

Fall 2018 course information:
Undergrad: INDENG 135 - LEC 001, 28108
Grad Section: INDENG 290 – Lec 2, 17093
Location: TuTh 12:30 pm – 2:00 pm | Evans 10

Prerequisite: Students should have a working knowledge of Python, have completed a fundamental probability / statistics course, as well as have a basic understanding Linear Algebra.

Welcome to Applied Data Science with Venture Applications

Sutardja Center for Entrepreneurship & Technology
Industrial Engineering End Operations Research Department

- Teaching Team
  - Ikhlaq Sidhu, SCET Faculty Director / Chief Scientist
  - Alexander Fred-Ojala, afo@berkeley.edu (instructor)
  - Sumayah Rahman, rahmans@berkeley.edu (GSI)
  - Tanya Piplani, tanyapiplani@berkeley.edu (GSI)

- Data-X Lab Students also contribute and coordinate: Vanessa Salas, and others..

Course Philosophy

Make the Tools

Use State-of-the-Art Open Source Tools

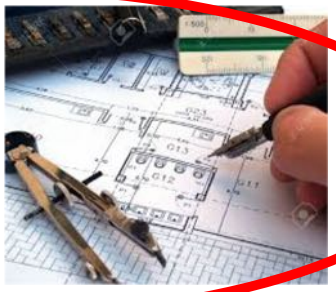

Architect the System

Sell, market, and pitch the product

Most CS / Math

Data-X

Business Topics

What is in this course

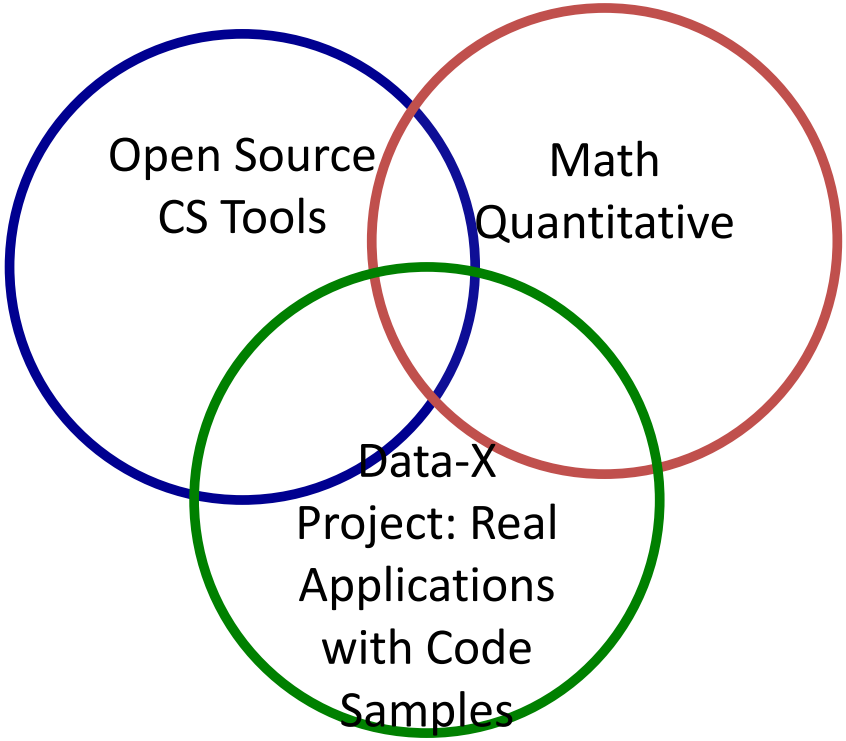

Holistic Perspective: Industry, Social Applications, Customer Driven

What is taught in the class?

- The ML stack most commonly used in creating ML/AI/Data applications
- Application and systems viewpoint of data and ML
- Implementation, architecture, and relevant processes to build real systems
- Connection with relevant mathematical, statistical foundations (optimization, entropy, correlation, LTI, prediction, classification)
- Practical insight into advanced techniques and tools: (eg. CNNs, NLP, scraping, recurrent networks, etc.)
- System modeling for data applications
- Application talks: Recommender systems, Blockchain, Spark etc.

Course Overview

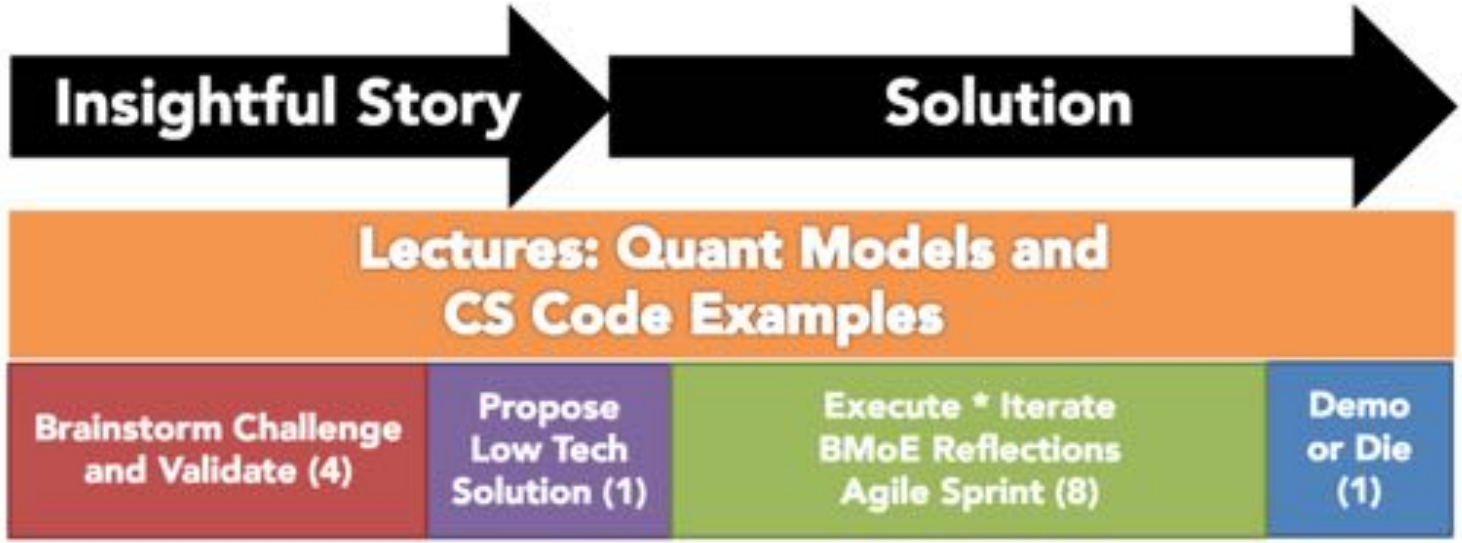

Open-ended, real-world project: Typically 5 students, with available advisor network

Data-X Project Examples

- Detection of fake news
- Prediction of long-term energy prices
- Automatic recycling through image recognition
- AI for crime detection, traffic guidance, medical diagnostics, etc.
- A version of Zillow that is recalculated with the effects of AirBnB income
- Signal processing and pattern analysis to improve earthquake warning systems
- Early Autism Detection
- Secure Health Records stored on a Blockchain

find many, many more at:

Data X www.data-x.blog/projects

Data-X By the numbers

Pantas and Ting
Sutardja Center
for Entrepreneurship & Technology
Berkeley Engineering

- 350 alumni students
- 50% avg enrollment increase / semester
- 80+ great projects completed
- 8+ published research papers
- 100+ industry experts in network
- 20+ students got employment as data scientists only because they took Data-X

Many Opportunities even After the course
Teaching
Data-Lab Projects
Develop New Data-X Materials

Amazing testimonials:

I think this class is so awesome because it teaches the tools and concepts that are most commonly used in workplace teams that are involved with data science machine learning.

Data X

Agenda on Day 1

Today:

- Course Introduction High Level Overview of Data, and Data-X Project (1:20 min)

By this week:

- Get your Notebook/development environment working
- Develop initial project ideas
- HW assigned by email/bcourse

Key Dates:

Class Starts today

Final projects: During end of class, possibly reading week

Top 2 projects will showcase at Collider Cup event

Most Resources Are Available at data-x.blog

1. Go to Data-X.blog
   - Syllabus
   - Instructions for SW Install
   - Link to GitHub with Cookbook Code Samples and Slides

2. Download Instructions to Install Python 3.x Anaconda Environment. For now you only need Anaconda, don't worry about other packages that are not already included.

3. Be able to create your own Jupyter notebook

4. Self-Review Python references as needed. See Ref CS01 and as needed BIDS Python Bootcamp.

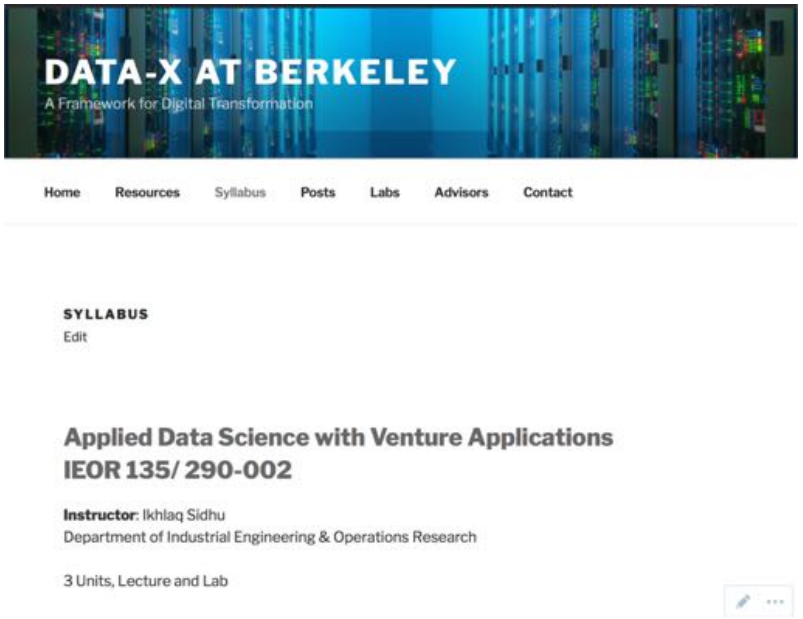

Project Types

Business or Consumer Use Case

Social Impact

(or improve part of a data pipeline or work towards a research result)

Its Just Cool

Homework For Week 1

- HW Part 1: For Your Project – By Next week
- Come up with 3 ideas for class projects in 1-3 sentences.
- A systems or application you will build
- **Communicate:** WHO the project is for, WHAT will it do, WHY this is needed/valuable.

- ---

- Homework Part II
- Python-based review notebook (Breakout and Homework, BKHW). To sent by email.

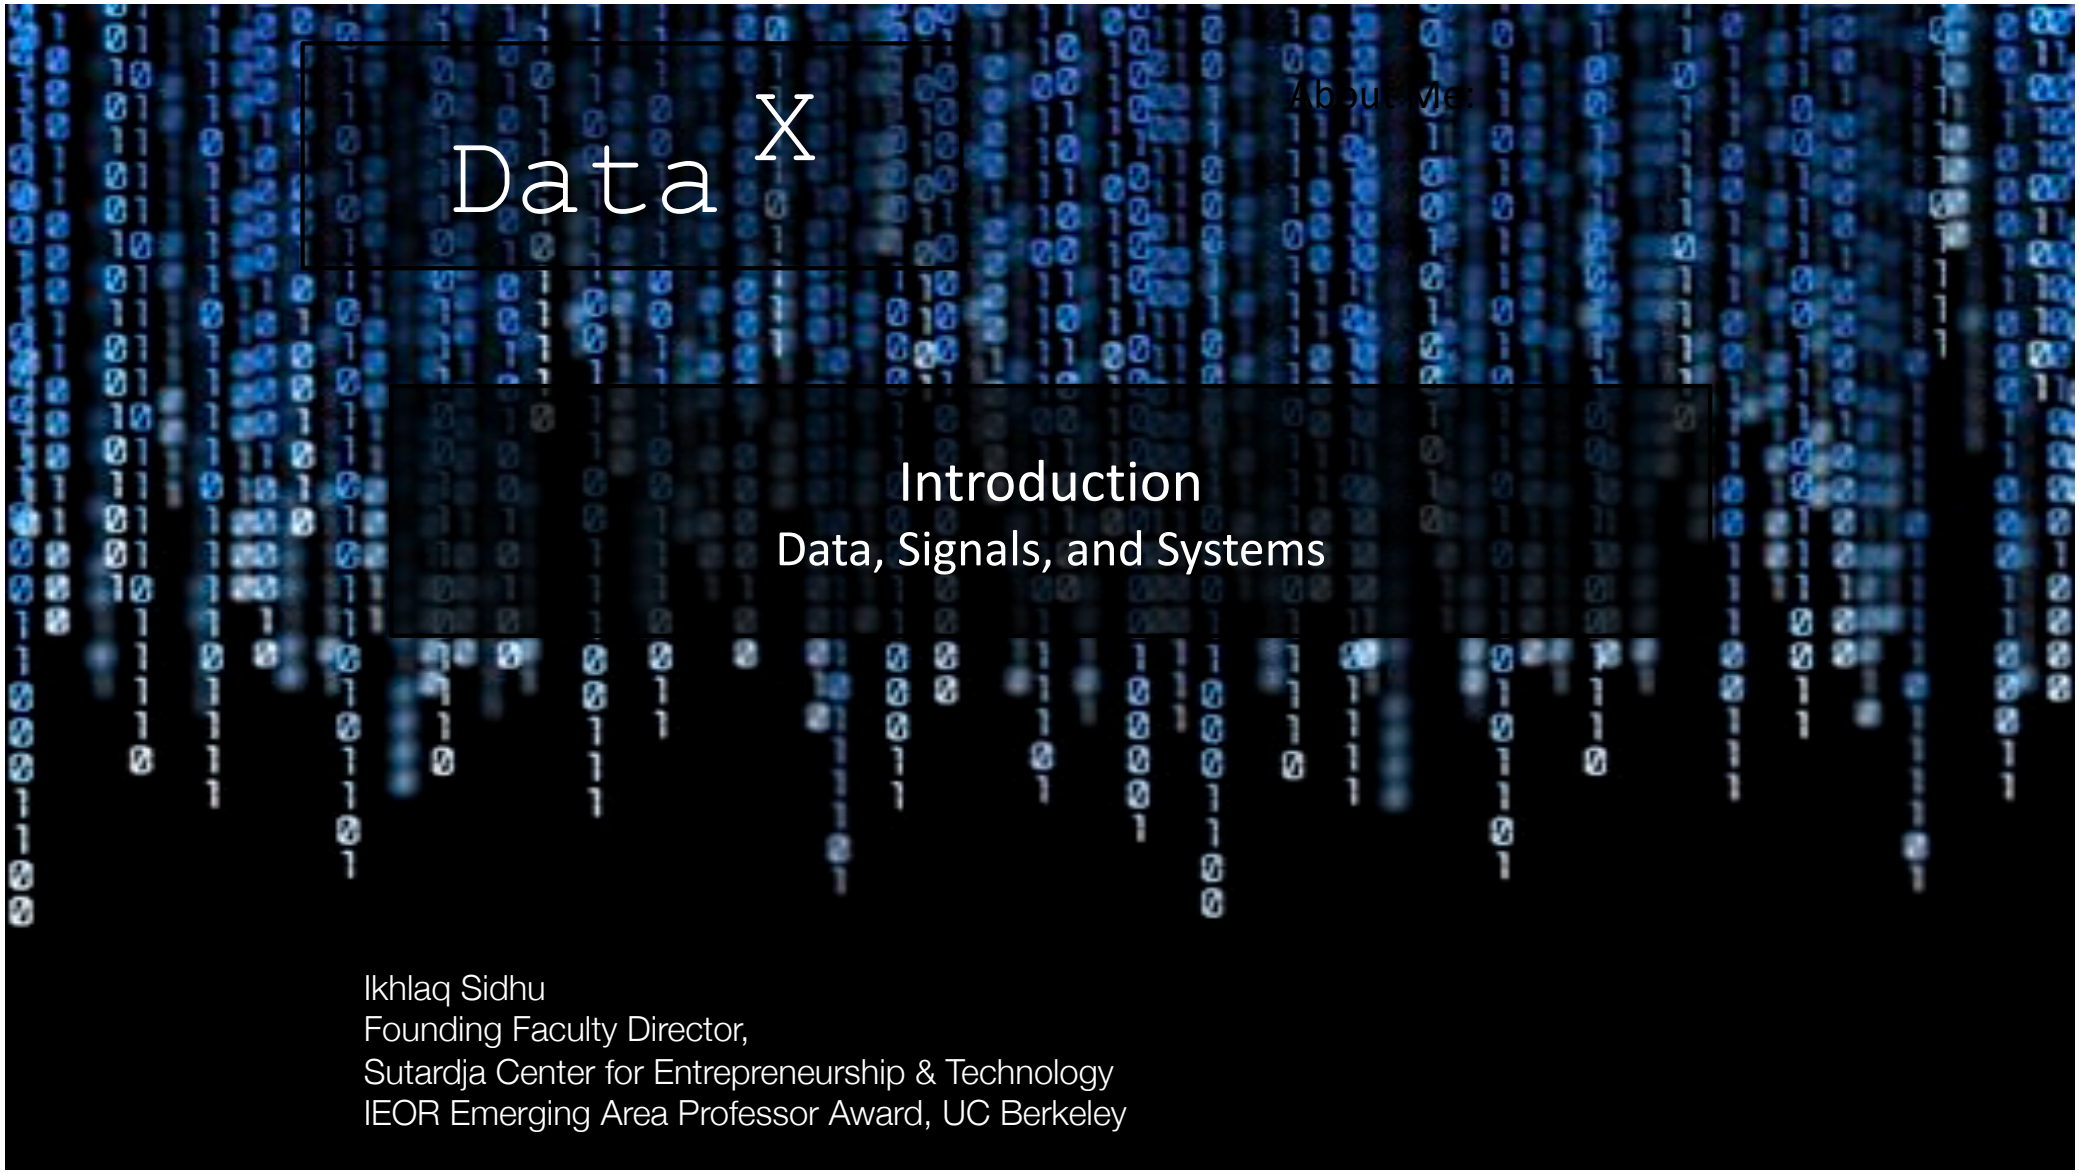
Data X
About Me:
Introduction
Data, Signals, and Systems
Ikhlaq Sidhu
Founding Faculty Director,
Sutardja Center for Entrepreneurship & Technology
IEOR Emerging Area Professor Award, UC Berkeley

An Overview of
Data and AI Applications

Basic Concept of Working with Data

- Data Wrangling
- In Production

Example: Data and Information is a competitive advantage

Real-life Example: ZestCash

- "All data is credit data"

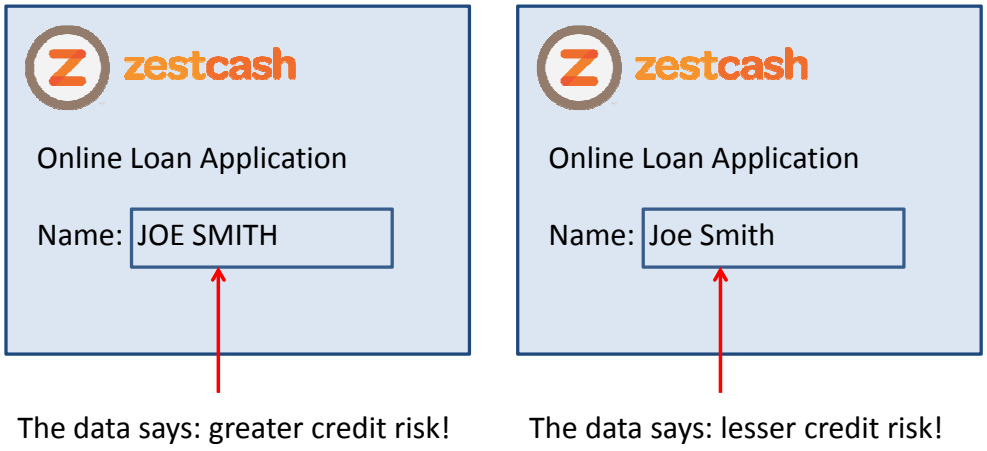

Reference: Shomit Ghose


• Service provider of Gambling and Casinos
• Entry Card
• Pain points
• Intervention
Harrah's Casino: Knowing your customer
PLAY & WIN
Reference: Supercrunchers

Why: More Simply

Customer Insight/ Engagement

Operations: Reliable & Predictable

Security & Fraud

TARGET

Harrah's

Walmart

Compliance 360

Financial Firms

Network Security

Implementation:
SW Tools / Stack

The Most Common Open Source Tools: AI/ML Stack

Start with Python as an interface
Jupyter Notebooks for prototyping

- Python: The interface
- NumPy, SciPy: Working with Arrays
- Pandas: Working in Tables, SQL to Pandas
- Sklearn: ML
- Matplotlib: Visualizing Data
- TensorFlow, Keras: Neural Networks
- SQL to Pandas
- NLP / NLTK: Natural Language
- Spark: For large data sets (GB, TB+)

https://www.youtube.com/watch?v=Q0jGAZAdZqM

https://conda.io/docs/user-guide/install/download.html

Where Does Data Come From?

Where Does Data Come From?

Real-life Example: ZestCash

- "All data is credit data"

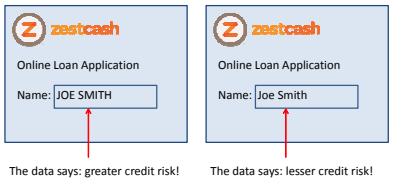

Your Own Web Site

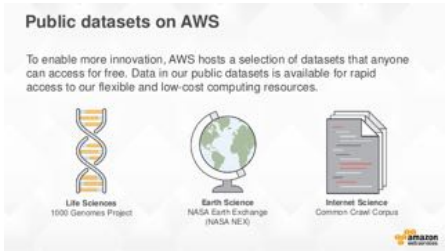

Public Data Sets
Stock market, etc.

IOT/Sensors

Other Web Sites

Web Scraping


```python
from bs4 import BeautifulSoup
import requests
page_link ='https://www.website_to_crawl.com'
# fetch the content from url
page_response = requests.get(page_link, timeout=5)
# parse html
page_content = BeautifulSoup(page_response.content, "html.parser")

# extract all html elements where price is stored
prices = page_content.find_all(class_='main_price')
# prices has a form:
#[<div class="main_price">Price: $66.68</div>,
# <div class="main_price">Price: $56.68</div>]

# you can also access the main_price class by specifying the tag of the class
prices = page_content.find_all('div', attrs={'class':'main_price'})
```

https://github.com/ikhlaqsidhu/data-x

https://github.com/ikhlaqsidhu/data-x/tree/master/03-tools-webscraping-crawling_api_afo

Formatting Data

An ML High Level Framework

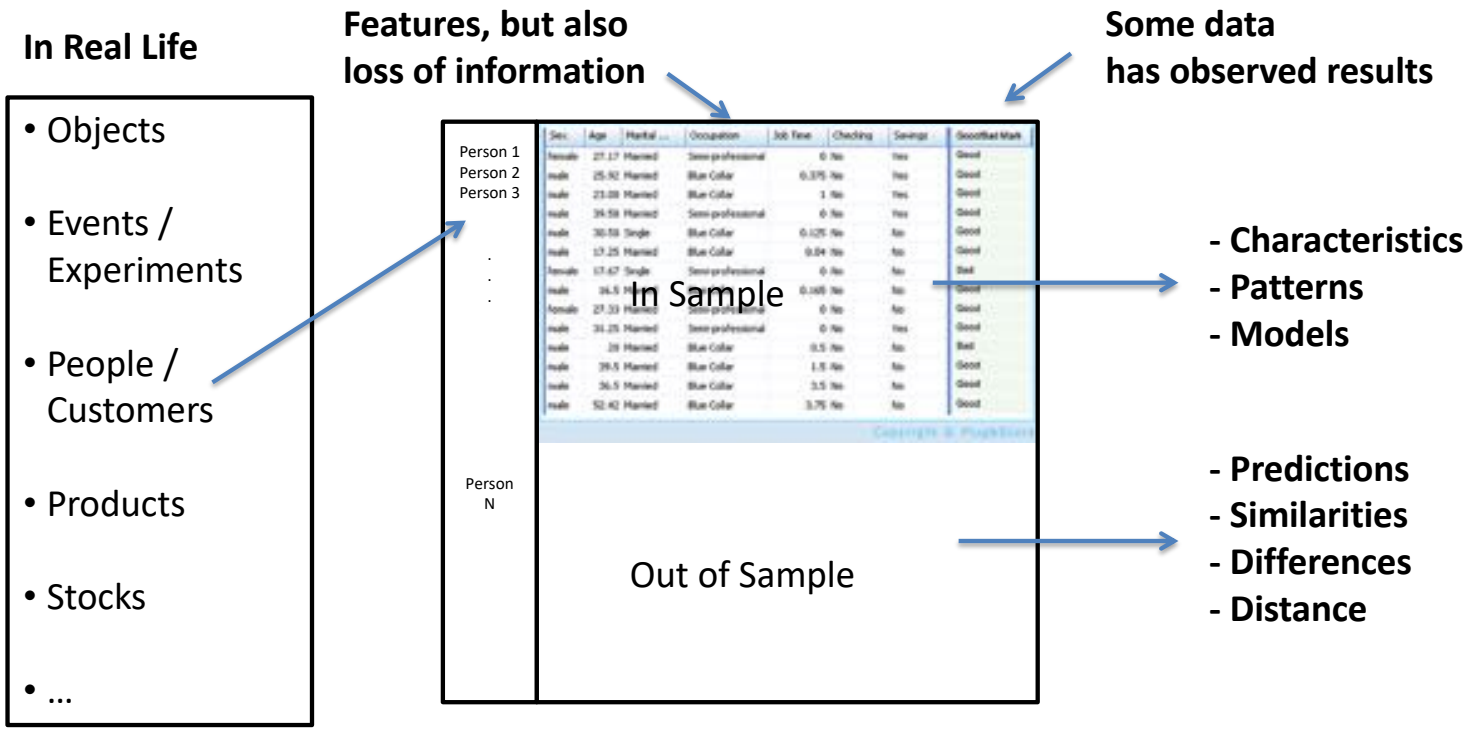

An ML High Level Framework

In Real Life

- Objects
- Events / Experiments
- People / Customers
- Products
- Stocks
- …

Features, but also loss of information

Some data has observed results

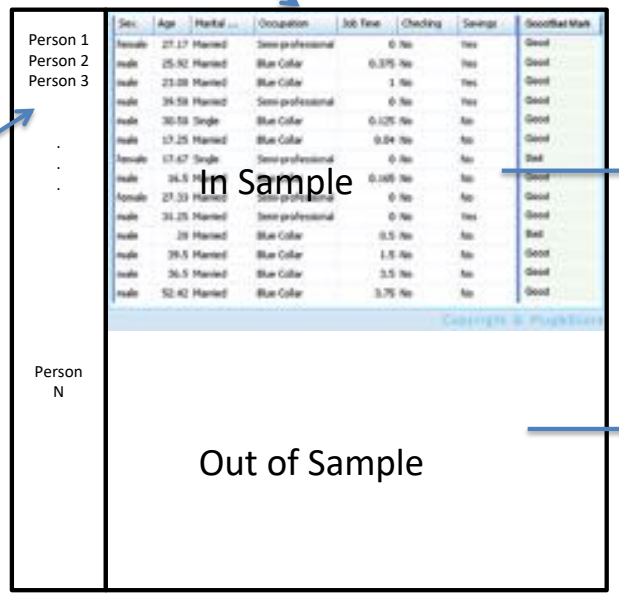

- Characteristics
- Patterns
- Models

- Predictions
- Similarities
- Differences
- Distance

CS: Table

Math: Matrix **X**, with N rows – each person
m columns, each feature (age, salary, ..)

$$\mathbf{X} = \begin{bmatrix} -2 & 4 & 7 & 31 \\ 6 & 9 & 12 & 6 \\ 12 & 11 & 0 & 1 \\ 9 & 10 & 2 & 3 \end{bmatrix}$$

A Fundamental Idea: From Table to Score

X =

Cust	F1	F2	F3
A	4	2	2
B	4.5	1.5	3
C	3	3	5
D	1	2	2
E	3	1.5	5
F	3.5	3.5	1
..	..	..	..

F(X)

Cust	Credit Score
A	552
B	381
C	760
D	330
E	452
F	678
..	..

A Fundamental Idea: From Table to Score

X =

Cust	F1	F2	F3
A	4	2	2
B	4.5	1.5	3
C	3	3	5
D	1	2	2
E	3	1.5	5
F	3.5	3.5	1
..	..	..	..

F(X)

Cust	Credit Score
A	552
B	381
C	760
D	330
E	452
F	678
..	..

```
#Setting up for Supervised learning
# First clean: use mapping + buckets

# X = matrix of data – e.g 1000 rows
# Y = In sample responses

# Typically we want to split in to training data and test data

X_train = X[0:500]
Y_train = Y[0:500]
X_test = X[501:1000]
Y_test = Y[501:1000]
```

Linear Regression Illustration

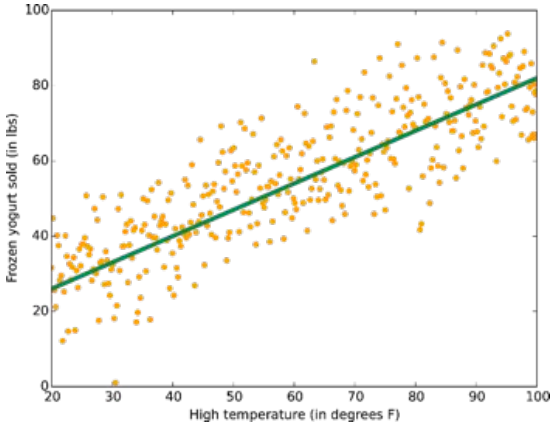


```
#Setting Linear Regression in sklearn
from sklearn import linear_model

model= linear_model.LinearRegression()
model.fit(X_train, Y_train)

Y_pred_train =  model.predict(X_train)
Y_pred_test =  model.predict(X_test)


# Compare Y_pred_test with Y_test for
error.
```

Illustration Source: https://docs.microsoft.com/en-us/azure/machine-learning/machine-learning-algorithm-choice

A Fundamental Idea: From Table to N- Dimensional Space

X =

Element	F1	F2	F3
A	4	2	2
B	4.5	1.5	3
C	3	3	5
D	1	2	2
E	3	1.5	5
F	3.5	3.5	1
..	..	..	..

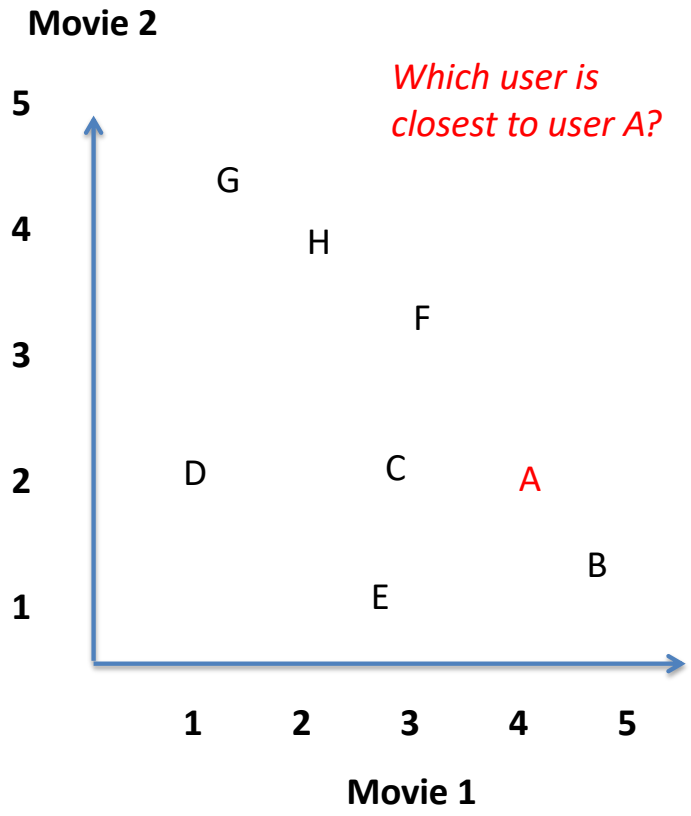

Clustering to Classification

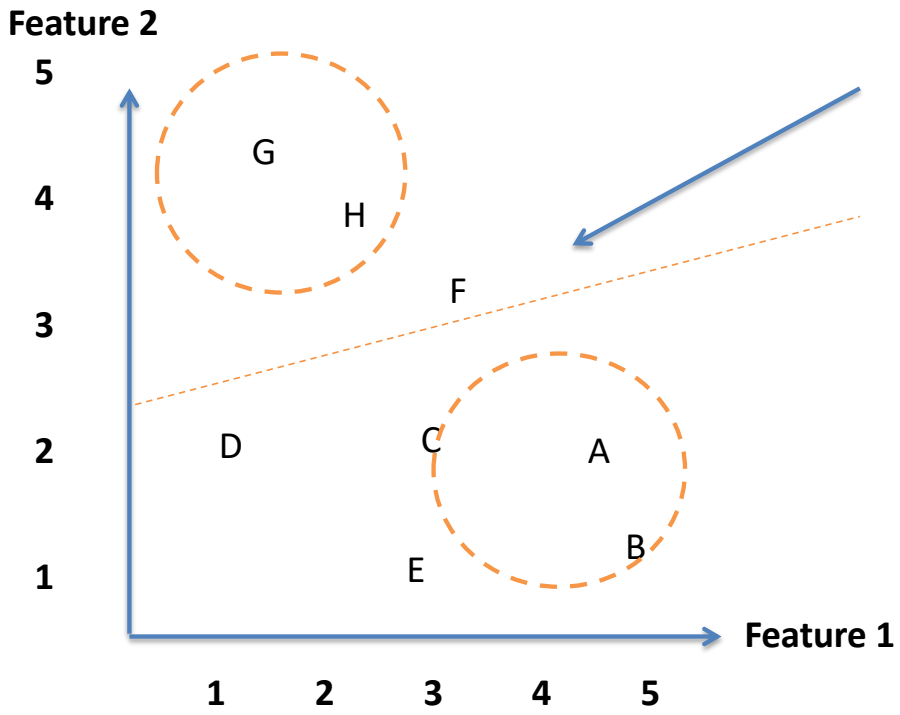

- Target customers?
- Pictures of Cats and Dogs
- Speech recognition
- Recognize Letters: A, B, C..

Traditionally 2 Tasks: Classification & Predictive Scoring

Extracted Data often in Table Format

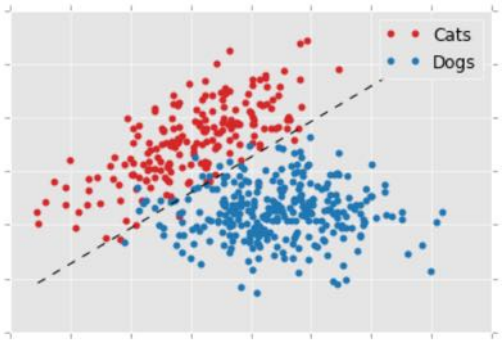

Classification:
Cats and Dogs, Speech Recognition
Movie Recommendation

Scoring:
Credit [illegible] Movie Rating
Heat [illegible]

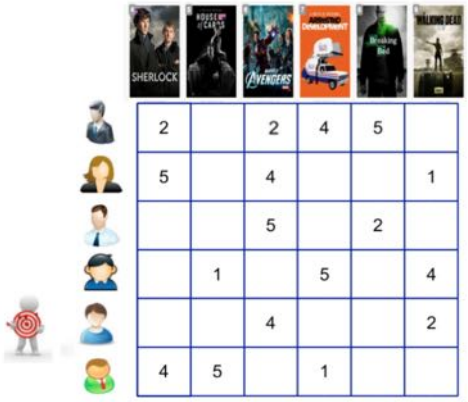

The most famous application has been recommendation: “which other user is most like you”

We have now switched to Neural Networks as Function Approximators

X → ML Algorithms Guess this function F(x) → Y

"Non-deep" feedforward neural network

X

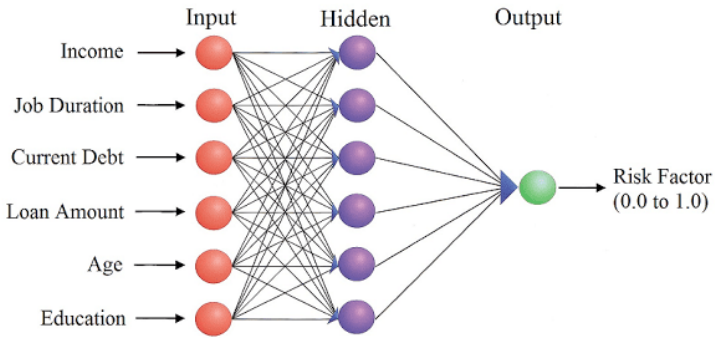

Y

Deep neural network

X

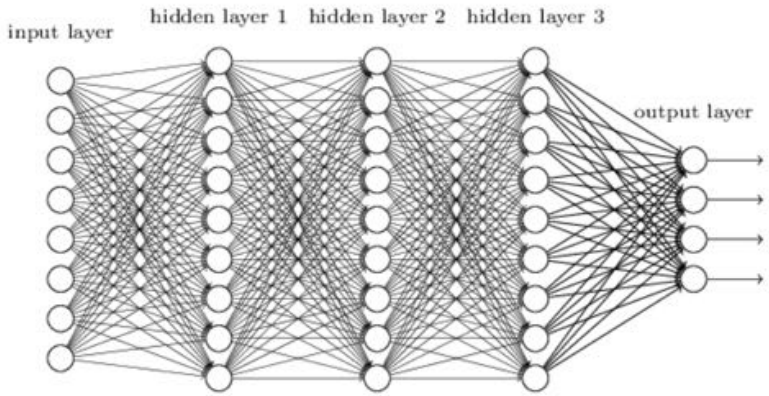

Y

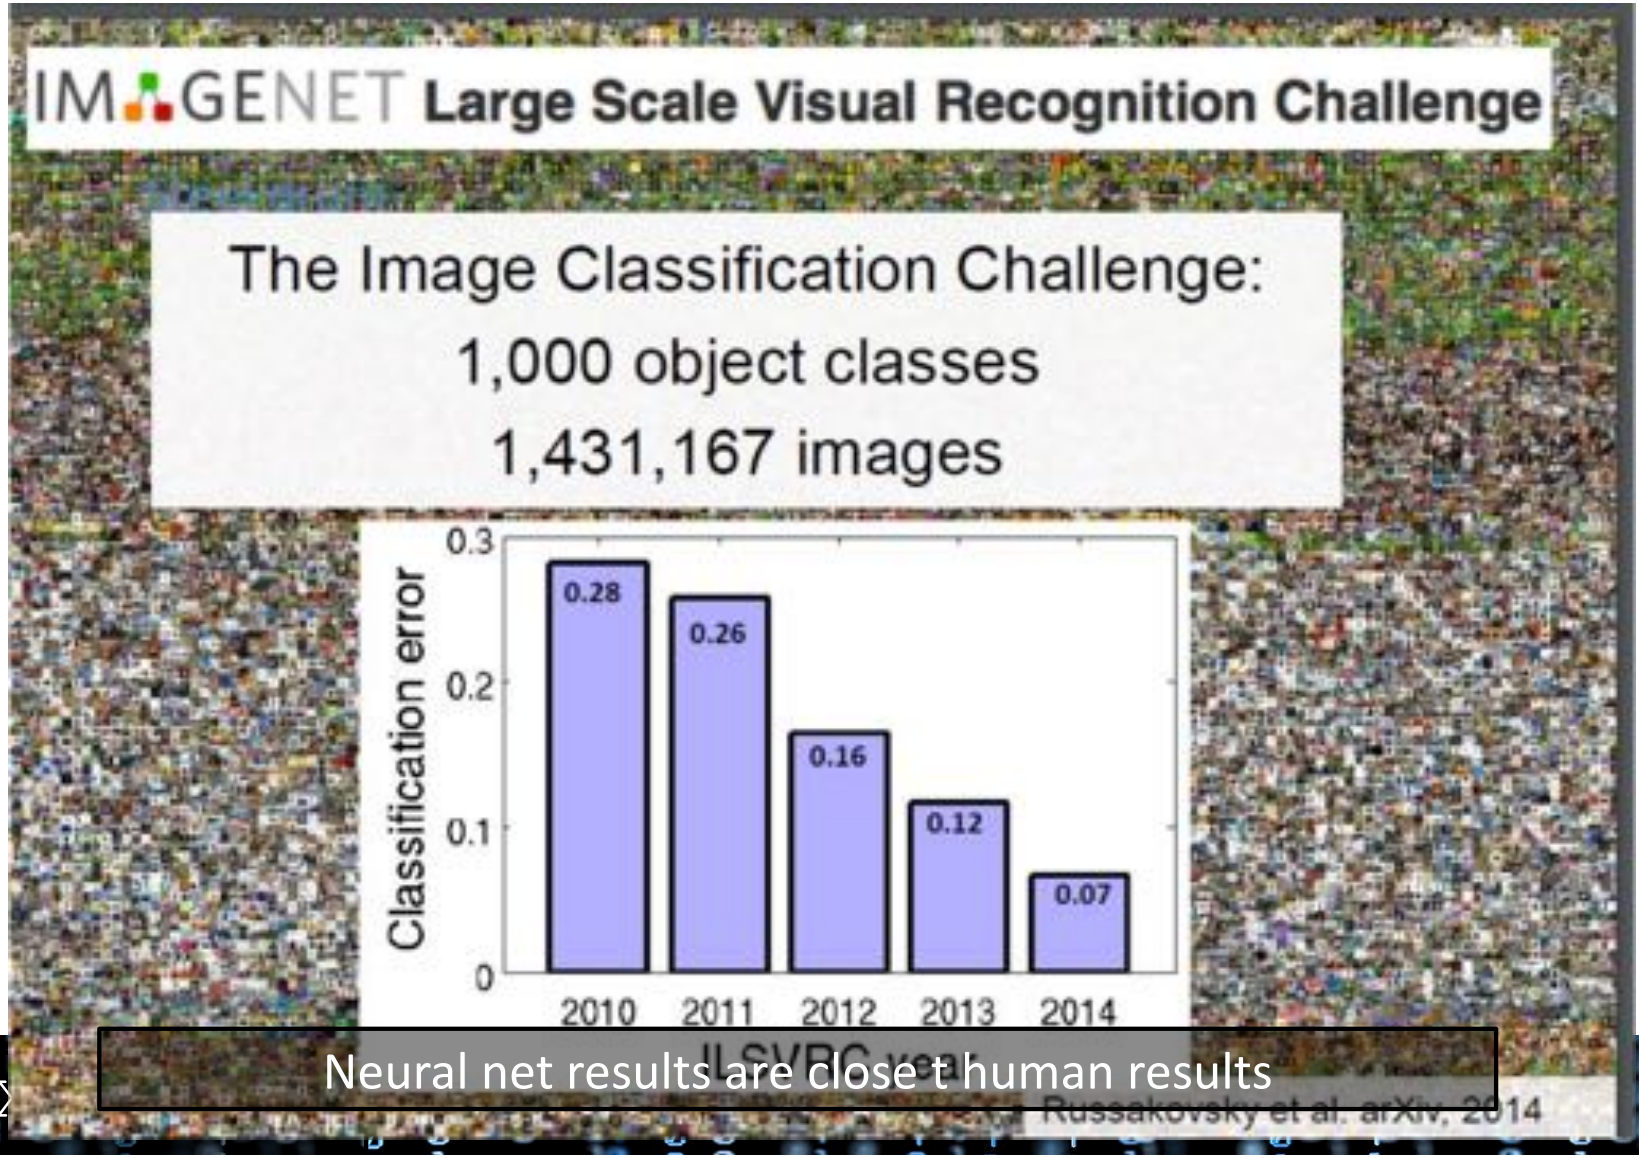

Neural net results are close t human results

Data

Project Types

Business or Consumer Use Case

Social Impact

(or improve part of a data pipeline or work towards a research result)

Its Just Cool

The Data-X System View

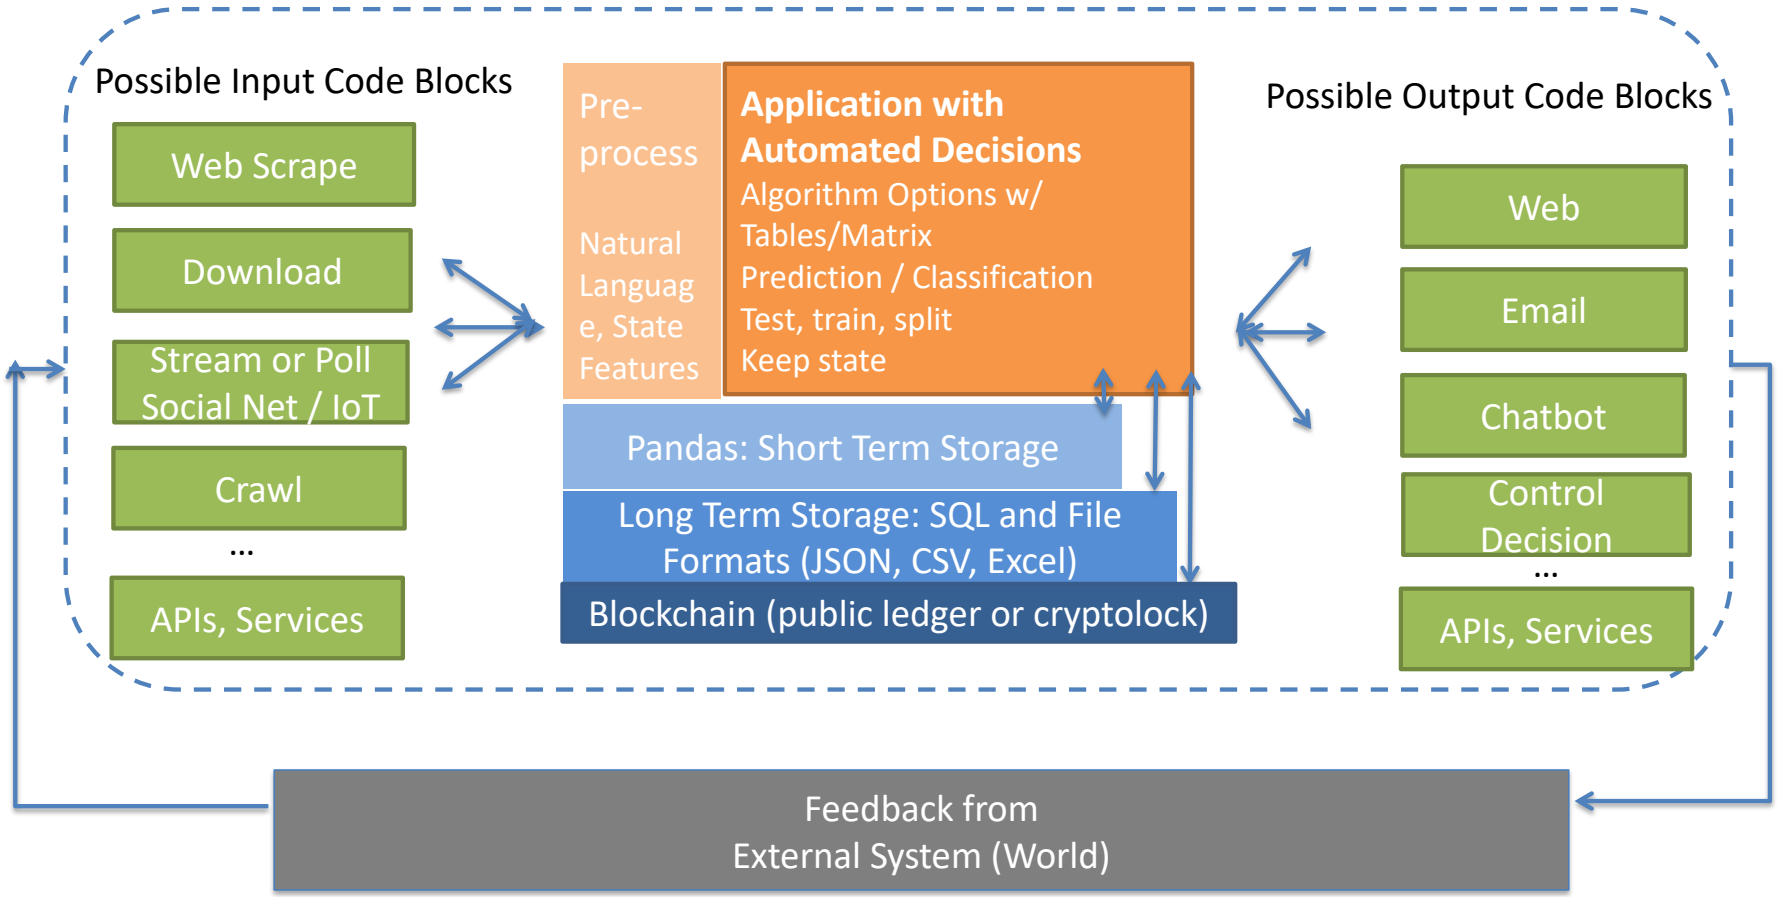

The Data-X System View: It's more than ML, it's also systems and models

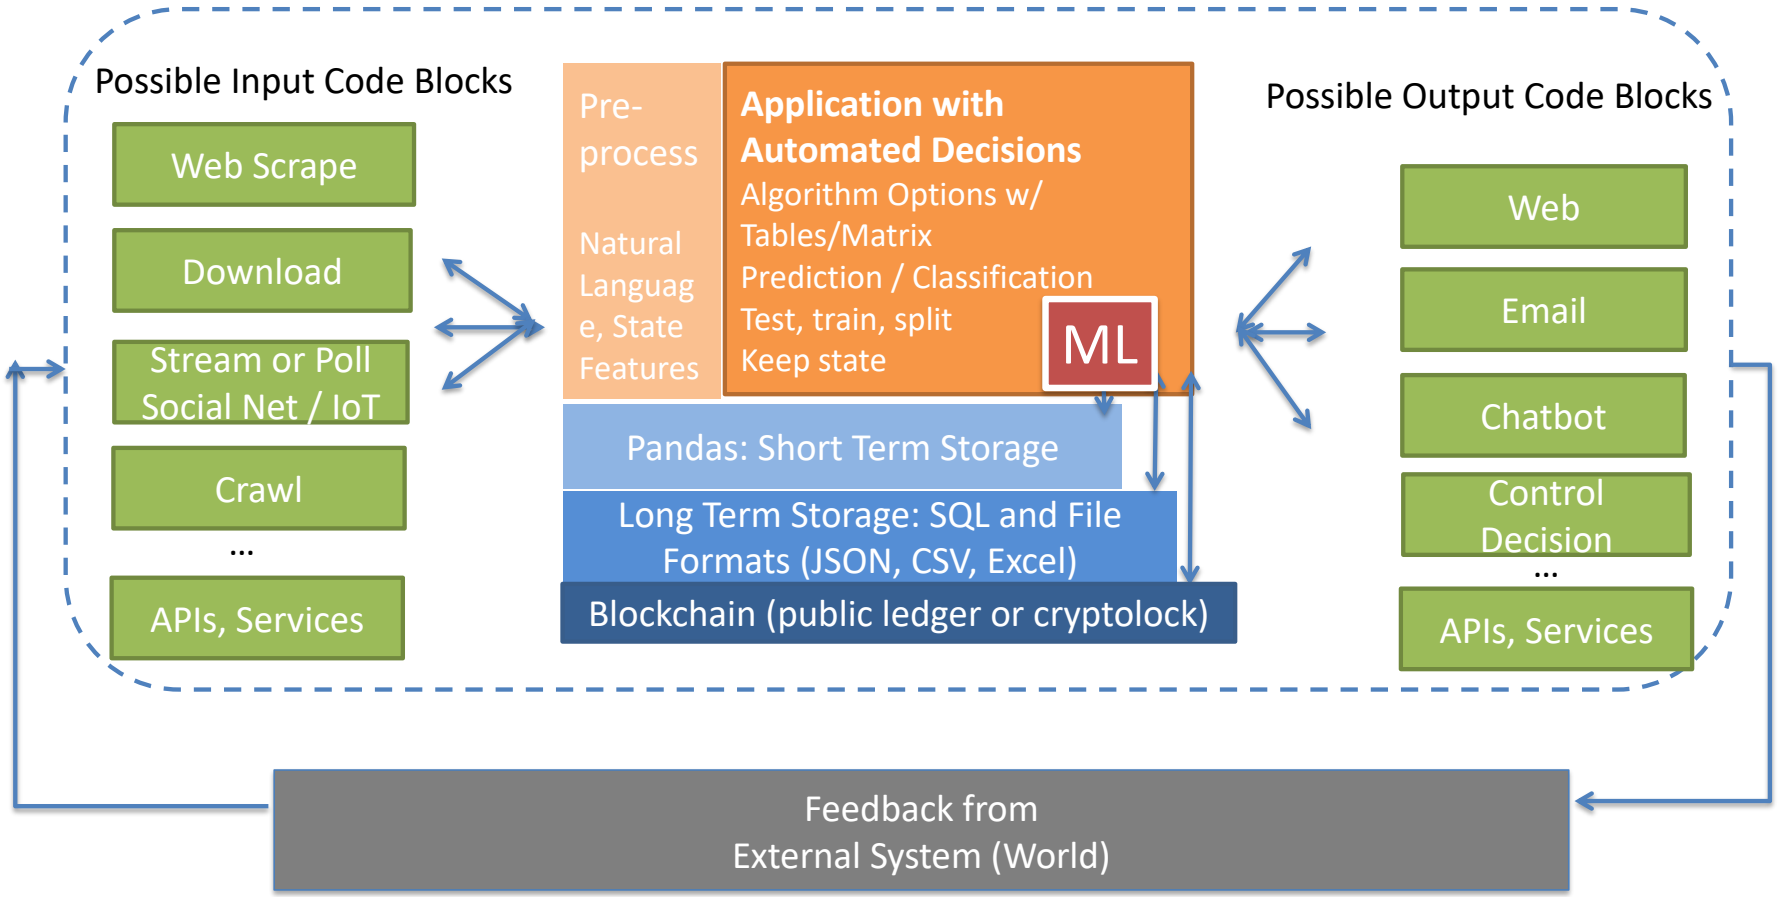

Project Ideation

- Past Projects Concepts:
  - See the Advisor's Tab of data-x.blog
- Past Projects:
  - See the archive on the Posts page and on the Labs page of Data-x.blog
- Combine ideas or extend previous work
- You can also choose to build part of a system,
  - Ie, just the part that automatically collects data by web scraping, or
  - just the part that makes a decision based on data already available

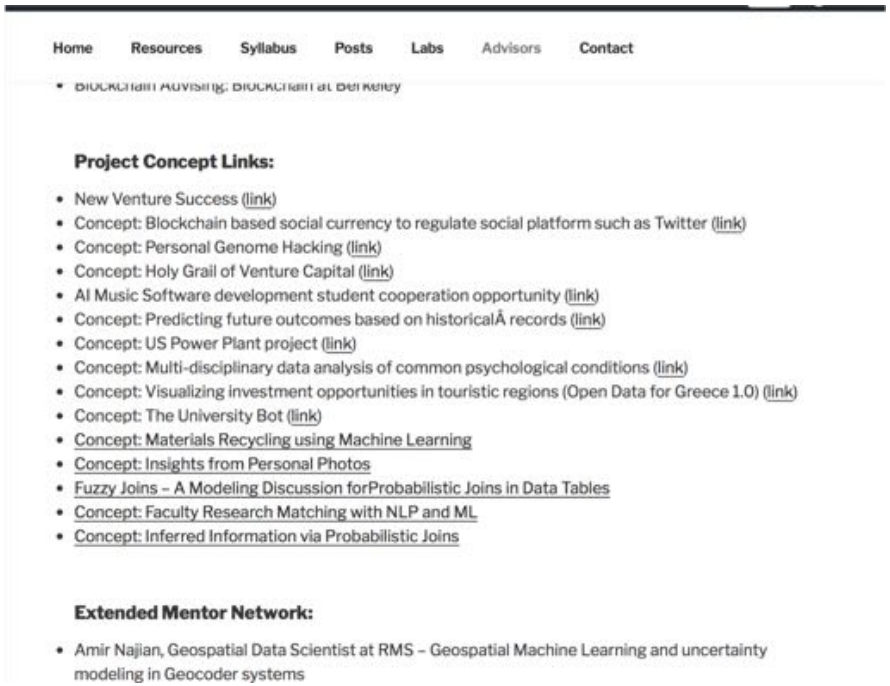

Home Resources Syllabus Posts Labs Advisors Contact

- Blockchain Advising: Blockchain at Berkeley

Project Concept Links:

- New Venture Success (link)
- Concept: Blockchain based social currency to regulate social platform such as Twitter (link)
- Concept: Personal Genome Hacking (link)
- Concept: Holy Grail of Venture Capital (link)
- AI Music Software development student cooperation opportunity (link)
- Concept: Predicting future outcomes based on historicalÃ records (link)
- Concept: US Power Plant project (link)
- Concept: Multi-disciplinary data analysis of common psychological conditions (link)
- Concept: Visualizing investment opportunities in touristic regions (Open Data for Greece 1.0) (link)
- Concept: The University Bot (link)
- Concept: Materials Recycling using Machine Learning
- Concept: Insights from Personal Photos
- Fuzzy Joins – A Modeling Discussion forProbabilistic Joins in Data Tables
- Concept: Faculty Research Matching with NLP and ML
- Concept: Inferred Information via Probabilistic Joins

Extended Mentor Network:

- Amir Najian, Geospatial Data Scientist at RMS – Geospatial Machine Learning and uncertainty modeling in Geocoder systems

Reminder: Homework For Week 1

- HW Part 1: For Your Project – By Next week
- Come up with 3 ideas for class projects in 1-3 sentences.
- A systems or application you will build
- **Communicate:** WHO the project is for, WHAT will it do, WHY this is needed/valuable.

- ---

- Homework Part II
- Python-based review notebook (Breakout and Homework, BKHW). To sent by email.

End of Section