



# Data X

About Me:

## Gradient Descent, Classification & Logistic Regression

Alexander Fred-Ojala, Ikhlaz Sidhu, Kevin Bozhe Li

# Data-X Fall 2018

## Lecture 7: Outline

1. Linear Regression recap
2. Gradient Descent
3. Feature scaling
4. Intro to Classification
5. Logistic Regression



## Recap: Linear Regression






# Recap: Prediction


Given some data:

$[(x_1, y_1), (x_2, y_2) \dots (x_m, y_m)]$

| x | y  |
|---|----|
| 2 | 10 |
| 4 | 5  |
| 3 | 9  |
| 5 | 4  |
| 6 | 3  |



**Objective:** Be able to predict y given new input x



# Recap: Simple Linear Regression


**Simple Linear Regression:** Hypothesis function  $h_{\theta}(x)$

$$\hat{y} = f(x, \theta) = h_{\theta}(x) = \theta^T x = \theta_0 + \theta_1 x_1$$

$$x = \begin{bmatrix} 1 \\ x_1 \end{bmatrix} \quad \text{x is given input}$$

**Objective:** fit the best linear function to the training data, i.e. find optimal parameters  $\theta$

$$\theta = \begin{bmatrix} \theta_0 \\ \theta_1 \end{bmatrix}$$



Data X

# Recap: Multiple Linear Regression

**Multiple Linear Regression:**  $\hat{y} = h_{\theta}(x) = \theta_0 + \theta_1 x_1 + \theta_2 x_2 + \dots + \theta_n x_n = \theta^T X$

$\theta = \begin{bmatrix} \theta_0 \\ \theta_1 \\ \vdots \\ \theta_n \end{bmatrix}$  is the parameter vector and

$X = \begin{bmatrix} x_0^{(1)} & x_1^{(1)} & \dots & x_n^{(1)} \\ x_0^{(2)} & x_1^{(2)} & \dots & x_n^{(2)} \\ \vdots & \vdots & \ddots & \vdots \\ x_0^{(m)} & x_1^{(m)} & \dots & x_n^{(m)} \end{bmatrix}$  is the feature vector and


$h_{\theta}(X) = \begin{bmatrix} h_{\theta}(x^{(1)}) \\ h_{\theta}(x^{(2)}) \\ \vdots \\ h_{\theta}(x^{(m)}) \end{bmatrix}$  is the hypotheses vector

**Example of multiple linear regression (2 features)**

$x_1$  = TV advertising

$x_2$  = Radio advertising

$y$  = Sales



# Recap: Cost function (MSE)


## Simple Linear Regression

$$\hat{y} = h_{\theta}(x) = \theta_0 + \theta_1 x_1$$

### Cost function:

Measures how good our predictions are (MSE)

$$J(\theta) = \frac{1}{m} \sum_{i=1}^m (h_{\theta}(x^{(i)}) - y^{(i)})^2$$



# Recap: Minimize cost function

Optimal model parameters minimizes  $J(\theta)$

$$J(\theta) = \frac{1}{m} \sum_{i=1}^m (h_{\theta}(x^{(i)}) - y^{(i)})^2$$

$$J(\theta) = \frac{1}{m} (X\theta - y)^T (X\theta - y)$$

$$\min_{\theta} J(\theta) \implies \nabla_{\theta} J(\theta) = 0$$

$$\frac{\partial J}{\partial \theta} = 2X^T X\theta - 2X^T y = 0$$

Minimize by taking the [derivative w.r.t.  $\theta$ ] = 0

## Normal equation for Linear Regression

Closed form, analytical solution.


$$\theta = (X^T X)^{-1} X^T y$$

### Pros:

- Finds optimum in one calculation
- Really quick for small data sets

### Cons:

- $O(n^2 m)$  complexity, to calculate  $(X^T X)^{-1}$  can be slow if  $X$  has many features  $n$ .
- $(X^T X)^{-1}$  might not be invertible, ie singular (can be solved by using the pseduoinverse)






# Gradient Descent



# Introducing Gradient Descent

Gradient descent is an iterative optimization algorithm for finding the minimum of a function.

To reach minima one takes steps proportional to the negative gradient (or approximate gradient) of the function at the current point.



# Introducing Gradient Descent


Alternative way of minimizing the cost function:

$$J(\theta) = \frac{1}{m} (X\theta - y)^T (X\theta - y)$$

- *Will always converge because  $J(\theta)$  is convex*
- Start with / initialize  $\theta_0, \theta_1$  . E.g.  $(\theta_0, \theta_1) = (0, 0)$
- Keep changing  $\theta_0, \theta_1$  to reduce  $J(\theta_0, \theta_1)$ ,

## Illustration of Gradient Descent

for one parameter  $\theta$



Source: <https://sebastianraschka.com>

# Gradient Descent Algorithm: Linear Regression

1. Calculate the partial derivative  $\frac{\partial}{\partial \theta_j} J(\theta)$  for all  $j$
2. Form the **update rule** for every parameter:  
$$\theta_{j,iter+1} := \theta_{j,iter} - \alpha \frac{\partial}{\partial \theta_j} J(\theta) = \theta_{j,iter} - \alpha/m \sum_{i=1}^m (h_{\theta}(x^{(i)}) - y^{(i)})x_j^{(i)}$$
3. Choose a step size/ *learning rate*  $\alpha$  (often between  $10^{-6}$  and  $10^2$  -- not too big, then divergence).
4. Update all the parameters  $\theta_0 \dots \theta_n$  by feeding in all training samples in  $X$  ("batch" Gradient descent)
5. Stop when the error has converged.


$$J_{train}(\theta) = \frac{1}{2m} \sum_{i=1}^m (h_{\theta}(x^{(i)}) - y^{(i)})^2$$

Repeat {

$$\theta_j := \theta_j - \alpha \frac{1}{m} \sum_{i=1}^m (h_{\theta}(x^{(i)}) - y^{(i)})x_j^{(i)}$$

(for every  $j = 0, \dots, n$ )

}



Source:Ritchie Ng



# Gradient Descent Tips

## Feature Scaling:

Gradient Descent will be quicker and more stable if the features are scaled.

### Standardization

For all features:

- Subtract mean
- Divide by st.dev.



$$x_i \leftarrow \frac{x_i - \mu(x_i)}{\sigma(x_i)}$$

### Min-max scaling

For all features:

- Subtract min(x\_i)
- max(x\_i)-min(x\_i)


$$x_i \leftarrow \frac{x_i - \min(x_i)}{\max(x_i) - \min(x_i)}$$



# Gradient Descent Tips

## Monitor convergence

**Plot the value of the error function  $J(\theta)$  at every iteration.**  
*Check that the error becomes smaller. Plot for different learning rates to find the best one.*



# Gradient Descent

## Pros

- **Will always converge** if learning rate  $\alpha$  is chosen correctly
- **Fast** (time complexity is  $O(m)$ )
- Supports out of sample training (stochastic / mini batch G.D.)

## Cons

- **We have to choose learning rate** and initialize model parameters
- **Often takes many iterations**



# Classification








# Regression vs. Classification

## Regression:


- Continuous output  $y$
- Quantitative approach
- Linear or Non-linear




## Classification:

- Discrete output  $y$
- Qualitative approach
- Linear or Non-linear


Ex. KNN,  
Logistic, SVM, ..



**KNN Method:** Find the  $k$  nearest images and have them vote on the label (i.e. take the mode)




5-NN classifier




# Examples of classification

## Examples

- Weather: Sunny / Rainy
- Spam Detection
- Image Classification: Cats VS Dogs
- Image Classification: Recognizing Digits



80322-4129 80206  
40004 14310  
37879 05453  
3502 75216  
35460 14209



# Our Goal: Classify items

*i.e. find the best hypothesis function  $h_{\theta}(x)$  that maps  $x$  to  $y$*



$x_i$



**Model**

$$\hat{y} = f(x, \theta) = h_{\theta}(x)$$

## **Binary classification (cat vs dog):**

$y = 1$  if picture is dog

$y = 0$  if picture is cat

$$y \in \{0, 1\}$$

**We have this data:**

**(X,Y):**  $(x_1, y_1), (x_2, y_2) \dots (x_n, y_n)$

- $x_i$  and  $y_i$  are arrays for each data element
- **Example:**  $x_i = [12 \ 15] = [\text{height}, \text{weight}]$ ,  $y_i = \text{male} / \text{female}$
- **For a picture:**  $x_i = [32 \times 32 \times 3]$  multidim array,  $y_i = \text{cat} / \text{dog}$

## **Multi-class classification:**

$$y_i = [y_{i,0}, y_{i,1}, \dots, y_{i,k}]$$

$$y_i = [1, 0, \dots, 0] \quad y \in \{0, 1, 2, \dots, k\}$$

$y_{i,0} = 1$  if picture is a dog

$y_{i,1} = 1$  if picture is a cat

$y_{i,2} = 1$  if picture is an elephant


etc.

# Our Goal: To classify items.

We have this: (X,Y)



$x_i$



**Actual Results:**

$$y_i = [y_{i,1}, y_{i,2}, \dots, y_{i,k}]$$
$$y_i = [0, 1, \dots, 0]$$

Machine Learning Steps to **train** a classifier model

1. Choose **model**:  $h_{\theta}(x)$  = estimate of Y
2. Define a **loss function** ( $J(\theta)$ ) = which is a function of  $f(Y_{\text{actual}}, Y_{\text{estimated}})$
3. Optimize across the parameter space ( $\theta$ ) to **minimize the loss function**



# Linear Regression for Classification? (Not so good!)




## Why not choose a Linear model for classification?

Because a line is not a good estimator for binary results (classification)

Linear model:  $f(x, \theta) = h_{\theta}(x) = \theta x_i = \theta_0 + \theta_1 x_{i,1} + \theta_2 x_{i,2} \dots$

Negative probabilities and biased towards majority class

**Instead we use Logistic Regression!**



# Logistic Regression




# Classification Example

**Data: students study for an exam**  
(x= hours studied, y = pass/not pass)

| Hours | 0.50 | 0.75 | 1.00 | 1.25 | 1.50 | 1.75 | 1.75 | 2.00 | 2.25 | 2.50 | 2.75 | 3.00 | 3.25 | 3.50 | 4.00 | 4.25 | 4.50 | 4.75 | 5.00 | 5.50 |
|-------|------|------|------|------|------|------|------|------|------|------|------|------|------|------|------|------|------|------|------|------|
| Pass  | 0    | 0    | 0    | 0    | 0    | 0    | 1    | 0    | 1    | 0    | 1    | 0    | 1    | 0    | 1    | 1    | 1    | 1    | 1    | 1    |

**y, binary output**  
0 = fail  
1 = pass



## Problem:

We want to find a model that can predict the probability that the student passes given x hours of study

$$h_{\theta}(x) = P(y = 1|x; \theta)$$

If Prob  $\geq 0.5$ , predict student will pass,  $y=1$   
If Prob  $< 0.5$ , predict student will fail,  $y=0$

$$P(y = 0|x; \theta) = 1 - P(y = 1|x; \theta)$$

Data X

# The logistic / sigmoid function

The sigmoid function:

$$z(t) = \frac{e^t}{e^t + 1} = \frac{1}{1 + e^{-t}}$$


Large positive  $t$

$z \rightarrow 1$

Large negative  $t$

$z \rightarrow 0$

*This function only evaluates to values between 0 and 1 for all real numbers (like a probability)*



$t$ , sum of weighted inputs + bias

$$t = \theta_0 + x_1 \theta_1$$

$$z(t) = z(\theta^T x) = \frac{1}{1 + e^{-(\theta_0 + \theta_1 x_1)}} = h_\theta(x)$$

If  $\theta_1$  is small  $\rightarrow$  slow rise  
If  $\theta_1$  is large  $\rightarrow$  fast rise

- $z(\theta x)$  is the probability that  $y = 1$  given any  $x$
- The **decision boundary**, where the probability = 50%
- $z(\theta x) = \frac{1}{2}$  when  $e^{-(\theta_0 + x_1 \theta_1)} = 1$ , ie  $\theta_0 + x_1 \theta_1 = 0$



# Decision Boundary

The decision boundary separates our predicted categories from one another, in the feature space.

If we have two inputs,  $x_1$  and  $x_2$ , the decision boundary is the line when the predicted probability for  $y=0$  and  $y=1$  is equal to 50%

$$h_{\theta}(x) = z(\theta^T X) = \frac{1}{1 + e^{-(\theta_0 + \theta_1 x_1 + \theta_2 x_2)}} = 0.5$$

$\Leftrightarrow$

$$\theta_0 + \theta_1 x_1 + \theta_2 x_2 = 0$$


## Example

$$\theta_0 = -3 \quad \theta_1 = 1 \quad \theta_2 = 1$$

Then  $x_1 + x_2 - 3 \geq 0$

will predict  $y=1$  and vice versa

(see example below)



# Derivation of the logistic cost function



- How to choose parameters  $\theta$  to find the best  $h_{\theta}(x)$
- We need a **cost function  $J(\theta)$**  that measures performance, then find best  $\theta$ .
- Output  $y$  is binary and can only take on two values (0 or 1)
- Construct  **$J(\theta)$**  that penalizes wrong predictions

$$J(\theta) = \frac{1}{m} \sum_{i=1}^m \text{Cost}(h_{\theta}(x^{(i)}), y^{(i)})$$

$$\text{Cost}(h_{\theta}(x), y) = -\log(h_{\theta}(x)) \quad \text{if } y = 1$$

$$\text{Cost}(h_{\theta}(x), y) = -\log(1 - h_{\theta}(x)) \quad \text{if } y = 0$$

Cost plotted against  
predicted class probability  
when the true value is  
 $y=1$  (top) or  $y=0$  (bottom)



# Logistic cost function

**Cross Entropy for binary classification =**

$$Cost(h_{\theta}(x), y) = -y \log(h_{\theta}(x)) - (1 - y) \log(1 - h_{\theta}(x))$$

Actual output  $y$

Estimated output  $h_{\theta}(x)$

**Note:** Loss Function on the former slide can be added to form cross entropy for binary classification.

This cost function can be derived from the Maximum Likelihood estimation of the parameters.


# Gradient Descent & Logistic Regression

$J(\theta)$  = is a cost a function comparing our estimate  $h_{\theta}(x)$  and the true  $y$ .

Find optimal  $\theta$  (first initialize  $\theta$  with some random value)

Take small steps in the direction where  $J(\theta)$  is decreasing

**Update rule:**  $\theta_{j+1} = \theta_j - [(\text{step size } \alpha) \times \text{gradient of } J(\theta)]$



## Formal update rule

looks exactly like Linear Regression,  
but note that  $h_{\theta}(x)$  has changed)

$$J(\theta) = \frac{-1}{m} \left[ \sum_{i=1}^m y^{(i)} \log h_{\theta}(x^{(i)}) + (1 - y^{(i)}) \log (1 - h_{\theta}(x^{(i)})) \right]$$

Repeat {

$$\theta_j := \theta_j - \alpha \frac{\partial}{\partial \theta_j} J(\theta)$$

}

Same as:

Repeat {

$$\theta_j := \theta_j - \alpha \sum_{i=1}^m (h_{\theta}(x^{(i)}) - y^{(i)}) x_j^{(i)}$$

}



# Multi-class Logistic Regression: *One-vs-all* $y \in \{0, 1 \dots k\}$

## Sigmoid function:

$$z(t) = \frac{e^t}{e^t + 1} = \frac{1}{1 + e^{-t}}$$

Large  $t \rightarrow 1$ , Small  $t \rightarrow 0$


And if  $t$  has this form  
 $t = \theta x_i$  (in matrix form)  $= \theta_0 + \theta_1 x_{i,1} + \theta_2 x_{i,2} \dots$

$$z(t) = \frac{1}{1 + e^{-\theta^T x}} = \frac{1}{1 + e^{(-\theta_0 + \theta_1 x_1 + \theta_2 x_2 + \dots)}} = h_{\theta}(x)$$

- Easily extends to multiple features ( $x_1, x_2, x_3 \dots$ )
- And multiple parameter weights

## One-vs-all

- Take  $i$ :th class (against all other grouped into an alternative class), create decision boundary and calculate probability
- Final prediction will be the class that had the highest probability against all others.



$$h_{\theta}^{(0)}(x) = P(y = 0|x; \theta)$$

$$h_{\theta}^{(1)}(x) = P(y = 1|x; \theta)$$

...

$$h_{\theta}^{(k)}(x) = P(y = k|x; \theta)$$

$$\text{prediction} = \max_i (h_{\theta}^{(i)}(x))$$

y1 boundary  
y2 boundary  
y3 boundary



# Read about Softmax Regression for Multiclass Classification

p. 139 - 142 in the Textbook



End of Section



# References

- The material presented in this lecture references lecture material draws on the materials the following courses:
- UC Berkeley – CS 294-129 (Designing, Visualizing, and Understanding Deep Neural Networks):  
<https://bcourses.berkeley.edu/courses/1453965/pages/cs294-129-designing-visualizing-and-understanding-deep-neural-networks>
- Stanford – CS231n (Convolutional Neural Networks for Visual Recognition):  
<http://cs231n.stanford.edu/>
- Stanford – CS229 (Machine Learning) & Andrew Ng's Machine Learning at Coursera: <http://cs229.stanford.edu/> & <https://www.coursera.org/learn/machine-learning>




**Example Code:** Logistic Regression in Scikit-learn





# Example Code Sample with Logistic Regression Classifier



Input data

X: Attribute Information:

- sepal length in cm
- sepal width in cm
- petal length in cm
- petal width in cm

Y:

0 = 'setosa',  
1 = 'versicolor',  
2 = 'virginica'

```
print type (X)  
print X[0:5]
```

```
<type 'numpy.ndarray'>  
[[ 5.1  3.5  1.4  0.2]  
 [ 4.9  3.   1.4  0.2]  
 [ 4.7  3.2  1.3  0.2]  
 [ 4.6  3.1  1.5  0.2]  
 [ 5.   3.6  1.4  0.2]]
```

```
print Y[0:5]  
[0 0 0 0 0]
```

# Example Code Sample with Logistic Regression Classifier

```
1 → import numpy as np
    from sklearn import linear_model, datasets

    X = iris.data[:, 1:3] # only the first two features.
    Y = iris.target

    # https://en.wikipedia.org/wiki/Logistic_regression
2 → logreg = linear_model.LogisticRegression(C=1e5)

    # we create an instance of Neighbours Classifier and fit the data.
3 → logreg.fit(X, Y)

    # predict a category for every row in X
4 → Z = logreg.predict(X)
```

\* Z[2] will be the predicted number for row X[2]

**Class**  
**sklearn.linear\_model.**  
**LogisticRegression**

```
(penalty='l2',
 dual=False,
 tol=0.0001,
 C=1.0,
 fit_intercept=True,
 intercept_scaling=1,
 class_weight=None,
 random_state=None,
 solver='liblinear', max_iter=100,
 multi_class='ovr', verbose=0,
 warm_start=False,
 n_jobs=1)
```

[http://scikit-learn.org/stable/modules/generated/sklearn.linear\\_model.LogisticRegression.html](http://scikit-learn.org/stable/modules/generated/sklearn.linear_model.LogisticRegression.html)

Data X

# Code Samples with SciKit Learn

```
# Plot the decision boundary. For that, we will assign a color to each
# point in the mesh [x_min, x_max]x[y_min, y_max].
x_min, x_max = X[:, 0].min() - .5, X[:, 0].max() + .5
y_min, y_max = X[:, 1].min() - .5, X[:, 1].max() + .5
xx, yy = np.meshgrid(np.arange(x_min, x_max, h), np.arange(y_min, y_max, h))
Z = logreg.predict(np.c_[xx.ravel(), yy.ravel()])
# numpy.ravel: Return a contiguous flattened array.
```

xx shape is (171, 231)

yy shape is (171, 231)



np.c returns shape (39501, 2)

[[ 3.8 1.5 ]

[ 3.82 1.5 ]

[ 3.84 1.5 ] ...]

Z shape is (39501,)




# Plotting the Results

```
# Put the result into a color plot
Z = Z.reshape(xx.shape)
plt.figure(1, figsize=(4, 3))
plt.pcolormesh(xx, yy, Z, cmap=plt.cm.Paired)

# Plot also the training points
plt.scatter(X[:, 0], X[:, 1], c=Y, edgecolors='k',
            cmap=get_cmap("Spectral"))
plt.xlabel('Sepal length')
plt.ylabel('Sepal width')

# plt.xlim(xx.min(), xx.max())
# plt.ylim(yy.min(), yy.max())
# plt.xticks(())
# plt.yticks(())

plt.show()
```





# Methods for LogisticRegression

## Methods

|   |   |
|---|---|
| <code>decision_function(X)</code>         | Predict confidence scores for samples.  |
| <code>densify()</code>                    | Convert coefficient matrix to dense array format.   |
| <code>fit(X, y[, sample_weight])</code>   | Fit the model according to the given training data.   |
| <code>fit_transform(X[, y])</code>        | Fit to data, then transform it.   |
| <code>get_params([deep])</code>           | Get parameters for this estimator.  |
| <code>predict(X)</code>                   | Predict class labels for samples in X.  |
| <code>predict_log_proba(X)</code>         | Log of probability estimates.   |
| <code>predict_proba(X)</code>             | Probability estimates.  |
| <code>score(X, y[, sample_weight])</code> | Returns the mean accuracy on the given test data and labels.                                |
| <code>set_params(*params)</code>          | Set the parameters of this estimator.   |
| <code>sparsify()</code>                   | Convert coefficient matrix to sparse format.  |
| <code>transform(*args, **kwargs)</code>   | DEPRECATED: Support to use estimators as feature selectors will be removed in version 0.19. |



```
fit(X, y, sample_weight=None)
```

[\[source\]](#)

Fit the model according to the given training data.

**Parameters:** **X** : {array-like, sparse matrix}, shape (n\_samples, n\_features)

Training vector, where n\_samples is the number of samples and n\_features is the number of features.

**y** : array-like, shape (n\_samples,)

Target vector relative to X.

**sample\_weight** : array-like, shape (n\_samples,) optional

Array of weights that are assigned to individual samples. If not provided, then each sample is given unit weight.

*New in version 0.17: sample\_weight support to LogisticRegression.*

**Returns:** **self** : object

Returns self.

```
predict(X)
```

[\[source\]](#)

Predict class labels for samples in X.

**Parameters:** **X** : {array-like, sparse matrix}, shape = [n\_samples, n\_features]

Samples.

**Returns:** **C** : array, shape = [n\_samples]

Predicted class label per sample.

Fit and predict  
from ScikitLearn



# Regularization

**Why:** To avoid over-fitting

**How:** You penalize your loss function by adding a multiple of an L1 (LASSO) or an L2 (Ridge) norm of your weights vector  $w$

Your new loss function =  $L(X,Y) + \lambda N(w)$

**Tuning the regularization term  $\lambda$ :** Cross-validation:

- divide your training data,
- train your model for a fixed value of  $\lambda$  and test it on the remaining subsets
- repeat this procedure while varying  $\lambda$ .

Then you select the best  $\lambda$  that minimizes your loss function.



## Shrinkage Methods II: An example

