

Introduction to Neural Networks

Data-X Spring 2019

Tanya Piplani

What is a Neural Network?

- Like other machine learning methods that we saw earlier in the class , it is a technique to:

***map features** to labels or some dependent continuous value.*

or

- ***compute the function** that relates features to labels or some dependent continuous value.*

What is a Neural Network?

- Like other machine learning methods that we saw earlier in the class , it is a technique to:


*map **features** to labels or some dependent **continuous value**.*

or


$$f(y|x)$$


- **compute the function** that relates features to labels or some dependent continuous value.

Single (Biological) Neuron



[image source: cs231n.stanford.edu]

Single (Artificial) Neuron



Types of Learning

- Supervised Learning
 - Unsupervised Learning
 - Reinforcement Learning
-
- [also Transfer Learning, Imitation Learning , Meta Learning etc ...]

Supervised Learning: $X \rightarrow Y$

- Ex 1: Image Recognition
 - X = pixel values
 - Y = one hot vector encoding category



x



$$\begin{bmatrix} 1 \\ 0 \\ 0 \\ \vdots \\ 0 \end{bmatrix}$$

y



x




$$\begin{bmatrix} 0 \\ 1 \\ 0 \\ \vdots \\ 0 \end{bmatrix}$$

y

Unsupervised Learning


Unsupervised learning is a type of machine **Learning** algorithm used to draw inferences from datasets consisting of input data without labeled responses.



Neural Networks


- Origins: Algorithms that try to mimic the brain.
- Very widely used in 80s and early 90s; popularity diminished in late 90s.
- Recent resurgence: State-of-the-art technique for many applications
- Artificial neural networks are not nearly as complex or intricate as the actual brain structure

Neural networks




- Neural networks are made up of **nodes** or **units**, connected by **links**
- Each link has an associated **weight** and **activation level**
- Each node has an **input function** (typically summing over weighted inputs), an **activation function**, and an **output**

Neural networks



- Neural networks are made up of **nodes** or **units**, connected by **links**
- Each link has an associated **weight** and **activation level**
- Each node has an **input function** (typically summing over weighted inputs), an **activation function**, and an **output**


Neural networks



Fully Connected !!
Why ?

- Neural networks are made up of **nodes** or **units**, connected by **links**
- Each link has an associated **weight** and **activation level**
- Each node has an **input function** (typically summing over weighted inputs), an **activation function**, and an **output**


Neuron Model: Logistic Unit




Sigmoid (logistic) activation function: $g(z) = \frac{1}{1 + e^{-z}}$

Sigmoid Function

$$A = \frac{1}{1+e^{-x}}$$



Neuron Model: Logistic Unit




Why have non-linearity???

Sigmoid Function


$$A = \frac{1}{1+e^{-x}}$$

Sigmoid (logistic) activation function: $g(z) = \frac{1}{1 + e^{-z}}$



A feed-forward neural network with linear activation and any number of hidden layers is equivalent to just a linear neural network with no hidden layer. For example let us consider the neural network in figure with two hidden layers and no activation-


**Why have
non-linearity???**




Sigmoid Function

$$A = \frac{1}{1+e^{-x}}$$

```
y = h2 * W3 + b3
= (h1 * W2 + b2) * W3 + b3
= h1 * W2 * W3 + b2 * W3 + b3
= (x * W1 + b1) * W2 * W3 + b2 * W3 + b3
= x * W1 * W2 * W3 + b1 * W2 * W3 + b2 * W3 + b3
= x * W' + b'
```







**Why have
non-linearity???**

Sigmoid Function

$$A = \frac{1}{1+e^{-x}}$$




Neural Network



Feed-Forward Process

- Input layer units are set by some exterior function (think of these as **sensors**), which causes their output links to be **activated** at the specified level
- Working forward through the network, the **input function** of each unit is applied to compute the input value
 - Usually this is just the weighted sum of the activation on the links feeding into this node
- The **activation function** transforms this input function into a final value
 - Typically this is a **nonlinear** function, often a **sigmoid** function corresponding to the “threshold” of that node

Neural Network



$a_i^{(j)}$ = “activation” of unit i in layer j

$\Theta^{(j)}$ = weight matrix controlling function mapping from layer j to layer $j + 1$


$$a_1^{(2)} = g(\Theta_{10}^{(1)} x_0 + \Theta_{11}^{(1)} x_1 + \Theta_{12}^{(1)} x_2 + \Theta_{13}^{(1)} x_3)$$

$$a_2^{(2)} = g(\Theta_{20}^{(1)} x_0 + \Theta_{21}^{(1)} x_1 + \Theta_{22}^{(1)} x_2 + \Theta_{23}^{(1)} x_3)$$

$$a_3^{(2)} = g(\Theta_{30}^{(1)} x_0 + \Theta_{31}^{(1)} x_1 + \Theta_{32}^{(1)} x_2 + \Theta_{33}^{(1)} x_3)$$


$$h_{\Theta}(x) = a_1^{(3)} = g(\Theta_{10}^{(2)} a_0^{(2)} + \Theta_{11}^{(2)} a_1^{(2)} + \Theta_{12}^{(2)} a_2^{(2)} + \Theta_{13}^{(2)} a_3^{(2)})$$

If network has s_j units in layer j and s_{j+1} units in layer $j+1$,
then $\Theta^{(j)}$ has dimension $s_{j+1} \times (s_j + 1)$.



$$Y = x_1 * w_1 + x_2 * w_2 + x_3 * w_3 + \dots + x_n * w_n \text{ --linear regression}$$

Example:



$$Y = x_1 * w_1 + x_2 * w_2 + x_3 * w_3 + \dots + x_n * w_n \text{ --linear regression}$$


For sample 1:

x	6	5	3	1
w	0.3	0.2	-0.5	0
Y = ?				

For sample 2:

x	20	5	3	1
w	0.3	0.2	-0.5	0
Y = ?				

Example:



$$Y = x_1 * w_1 + x_2 * w_2 + x_3 * w_3 + \dots + x_n * w_n \text{ --linear regression}$$


For sample 1:

x	6	5	3	1
w	0.3	0.2	-0.5	0
$Y = \text{sum}(x * w) = 1.3$				

For sample 2:

x	20	5	3	1
w	0.3	0.2	-0.5	0
$Y = \text{sum}(x * w) = 5.5$				

Lets us apply a **threshold function** on the output:



$$f(t) = \begin{cases} t & \text{if } t < 3 \\ 0 & \text{otherwise} \end{cases}$$

For sample 1:

x	6	5	3	1
w	0.3	0.2	-0.5	0


$$Y = f(\text{sum}(x * w)) = f(1.3) = 1.3$$

For sample 2:

x	20	5	3	1
w	0.3	0.2	-0.5	0

$$Y = f(\text{sum}(x * w)) = f(5.5) = 0$$

Now, if we apply a **logistic/sigmoid function** on the output it will squeeze all the output between 0 and 1 :



$$Y = \text{Sigmoid}(x_1 * w_1 + x_2 * w_2 + \dots + x_n * w_n) \text{ --logistic regression}$$

Logistic/sigmoid function

$$S(x) = \frac{1}{1 + e^{-x}} = \frac{e^x}{e^x + 1}.$$


For sample 1:

x	6	5	3	1
w	0.3	0.2	-0.5	0
$Y = \sigma(\text{sum}(x * w)) = \sigma(1.3) = 0.78$				

For sample 2:

x	20	5	3	1
w	0.3	0.2	-0.5	0
$Y = \sigma(\text{sum}(x * w)) = \sigma(5.5) = 0.99$				

Now, if we apply a **logistic/sigmoid function** on the output it will squeeze all the output between 0 and 1 :



$$Y = \text{Sigmoid}(x_1 * w_1 + x_2 * w_2 + \dots + x_n * w_n) \text{ --logistic regression}$$

Logistic/sigmoid function

$$S(x) = \frac{1}{1 + e^{-x}} = \frac{e^x}{e^x + 1}$$


For sample 1:

x	6	5	3	1
w	0.3	0.2	-0.5	0
Y = $\sigma(\text{sum}(x * w)) = \sigma(1.3) = 0.78$				

For sample 2:

x	20	5	3	1
w	0.3	0.2	-0.5	0
Y = $\sigma(\text{sum}(x * w)) = \sigma(5.5) = 0.99$				

Now, if we apply a threshold on the **logistic/sigmoid** output it will set the final output as 0 or 1 :



Logistic/sigmoid function

$$S(x) = \frac{1}{1 + e^{-x}} = \frac{e^x}{e^x + 1}$$

$$f(t) = \begin{cases} 1 & \text{if } t > 0.6 \\ 0 & \text{otherwise} \end{cases}$$

For sample 1:

x	6	5	3	1
w	0.3	0.2	-0.5	0
Y = f(σ(sum(x * w))) = f(σ(1.3)) = f(0.78) = 1				

For sample 2:


x	20	5	3	1
w	0.3	0.2	-0.5	0
Y = f(σ(sum(x * w))) = f(σ(5.5)) = f(0.99) = 1				

Activation functions


We use **activation functions** in neurons to induce nonlinearity in the neural nets so that it can learn complex functions.

Sigmoid Function

$$A = \frac{1}{1+e^{-x}}$$



Tanh function




$$f(x) = \tanh(x) = \frac{2}{1+e^{-2x}} - 1$$

Problem -

- Towards either end of the sigmoid/tanh function, the Y values tend to respond very less to changes in X.
- The gradient at that region is going to be small. It gives rise to a problem of “vanishing gradients”


ReLU

$$f(x) = \max(0, x)$$



Leaky ReLU

$$f(x) = \begin{cases} 0.01x & \text{for } x < 0 \\ x & \text{for } x \geq 0 \end{cases}$$



Multiple Output Units: One-vs-Rest



Pedestrian




Car



Motorcycle



Truck



$$h_{\Theta}(\mathbf{x}) \in \mathbb{R}^K$$

We want:

$$h_{\Theta}(\mathbf{x}) \approx \begin{bmatrix} 1 \\ 0 \\ 0 \\ 0 \end{bmatrix}$$

when pedestrian

$$h_{\Theta}(\mathbf{x}) \approx \begin{bmatrix} 0 \\ 1 \\ 0 \\ 0 \end{bmatrix}$$

when car

$$h_{\Theta}(\mathbf{x}) \approx \begin{bmatrix} 0 \\ 0 \\ 1 \\ 0 \end{bmatrix}$$

when motorcycle

$$h_{\Theta}(\mathbf{x}) \approx \begin{bmatrix} 0 \\ 0 \\ 0 \\ 1 \end{bmatrix}$$


when truck

Representing Boolean Functions


Simple example: AND

$$x_1, x_2 \in \{0, 1\}$$

$$y = x_1 \text{ AND } x_2$$





$$h_{\theta}(\mathbf{x}) = g(-30 + 20x_1 + 20x_2)$$




x_1	x_2	$h_{\theta}(\mathbf{x})$
0	0	$g(-30) \approx 0$
0	1	$g(-10) \approx 0$
1	0	$g(-10) \approx 0$
1	1	$g(10) \approx 1$

Representing Boolean Functions



A OR B





Input		Output
A	B	$Y=A+B$
0	0	0
0	1	1
1	0	1
1	1	1



(NOT x_1) AND (NOT x_2)



????

Representing Boolean Functions



INPUTS		OUTPUT
X	Y	Z
0	0	1
0	1	0
1	0	0
1	1	0

Neural Network Learning



What does the machine learn?

One question that people often have when getting started in ML is:

“What does the machine (i.e. the statistical model) actually learn?”

This will vary from model to model, but in simple terms the model learns a function f such that $f(X)$ maps to y . Put differently, the model learns how to take X (i.e. features, or, more traditionally, independent variable(s)) in order to predict y (the target, response or more traditionally the dependent variable).

In the case of the simple linear regression ($y \sim \mathbf{b0} + \mathbf{b1} * X$ where X is one column/variable) the model “learns” (read: estimates) two parameters;

- $\mathbf{b0}$: the bias (or more traditionally the intercept); and,
- $\mathbf{b1}$: the slope

Learning parameters: Cost functions

Remember that in ML, the focus is on **learning from data**.
cost function—it helps the learner to correct / change behaviour to **minimize mistakes**.

In ML, cost functions are used to estimate how badly models are performing.
Put simply, *a cost function is a measure of how wrong the model is in terms of its ability to estimate the relationship between X and y .*


Example

$$h_{\theta}(x) = \theta_0 + \theta_1 x$$

The theta values are the *parameters*.

$$\theta_0 = 1.5$$
$$\theta_1 = 0$$


This yields $h(x) = 1.5 + 0x$. $0x$ means no slope, and y will always be the constant 1.5. This looks like:




How about

$$\theta_0 = 1$$

$$\theta_1 = 0.5$$



```
x = [1, 1, 2, 3, 4, 3, 4, 6, 4]
y = [2, 1, 0.5, 1, 3, 3, 2, 5, 4]
```



You want to
match your
hypothesis to
your best fit line

```
x = [1, 1, 2, 3, 4, 3, 4, 6, 4]  
y = [2, 1, 0.5, 1, 3, 3, 2, 5, 4]
```

$$\text{MSE} = \frac{1}{n} \sum_{i=1}^n (Y_i - \hat{Y}_i)^2.$$

Mean Squared Error

```
x = [1, 1, 2, 3, 4, 3, 4, 6, 4]  
y = [2, 1, 0.5, 1, 3, 3, 2, 5, 4]
```

$$\text{MSE} = \frac{1}{n} \sum_{i=1}^n (Y_i - \hat{Y}_i)^2.$$

$$= \frac{1}{m} \sum_{i=1}^m (h_{\theta}(x^{(i)}) - y^{(i)})^2$$

Backpropagation

Given the cost function how do you update parameters?

1. Calculate the **partial derivative** $\frac{\partial}{\partial \theta_j} J(\theta)$ for all j

2. Form the **update rule** for every parameter:

$$\theta_{j,iter+1} := \theta_{j,iter} - \alpha \frac{\partial}{\partial \theta_j} J(\theta) = \theta_{j,iter} - \alpha/m \sum_{i=1}^m (h_{\theta}(x^{(i)}) - y^{(i)})x_j^{(i)}$$

$$J_{train}(\theta) = \frac{1}{2m} \sum_{i=1}^m (h_{\theta}(x^{(i)}) - y^{(i)})^2$$

Repeat {

$$\theta_j := \theta_j - \alpha \frac{1}{m} \sum_{i=1}^m (h_{\theta}(x^{(i)}) - y^{(i)})x_j^{(i)}$$

(for every $j = 0, \dots, n$)

}

Motivation: loss minimization

Optimization problem:

$$\min_{\mathbf{V}, \mathbf{w}} \text{TrainLoss}(\mathbf{V}, \mathbf{w})$$

$$\text{TrainLoss}(\mathbf{V}, \mathbf{w}) = \frac{1}{|\mathcal{D}_{\text{train}}|} \sum_{(x, y) \in \mathcal{D}_{\text{train}}} \text{Loss}(x, y, \mathbf{V}, \mathbf{w})$$

$$\text{Loss}(x, y, \mathbf{V}, \mathbf{w}) = (y - f_{\mathbf{V}, \mathbf{w}}(x))^2$$

$$f_{\mathbf{V}, \mathbf{w}}(x) = \sum_{j=1}^k w_j \sigma(\mathbf{v}_j \cdot \phi(x))$$

Goal: compute gradient

$$\nabla_{\mathbf{V}, \mathbf{w}} \text{TrainLoss}(\mathbf{V}, \mathbf{w})$$

Approach


Mathematically: just grind through the chain rule

Next: visualize the computation using a computation graph

Advantages:

- Avoid long equations
- Reveal structure of computations (modularity, efficiency, dependencies)

Functions as boxes








$$2\text{in}_1 + \text{in}_2 * \text{in}_3 = \text{out}$$

Partial derivatives (gradients): how much does the output change if an input changes?


Example:

$$2\text{in}_1 + (\text{in}_2 + \epsilon)\text{in}_3 = \text{out} + \text{in}_3\epsilon$$

Basic building blocks



Composing functions




Chain rule:

$$\frac{\partial \text{out}}{\partial \text{in}} = \frac{\partial \text{out}}{\partial \text{mid}} \frac{\partial \text{mid}}{\partial \text{in}}$$

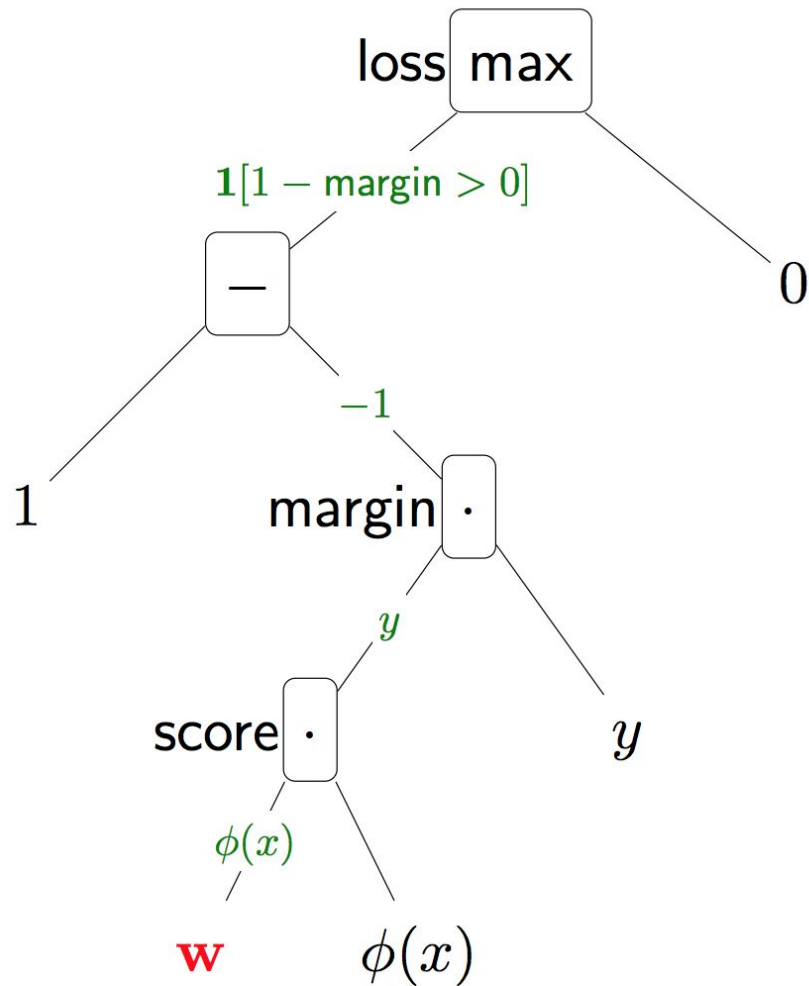
Example

$$\text{Loss}(x, y, \mathbf{w}) = \max\{1 - \mathbf{w} \cdot \phi(x)y, 0\}$$



Example

$$\text{Loss}(x, y, \mathbf{w}) = \max\{1 - \mathbf{w} \cdot \phi(x)y, 0\}$$




Gradient: multiply the edges

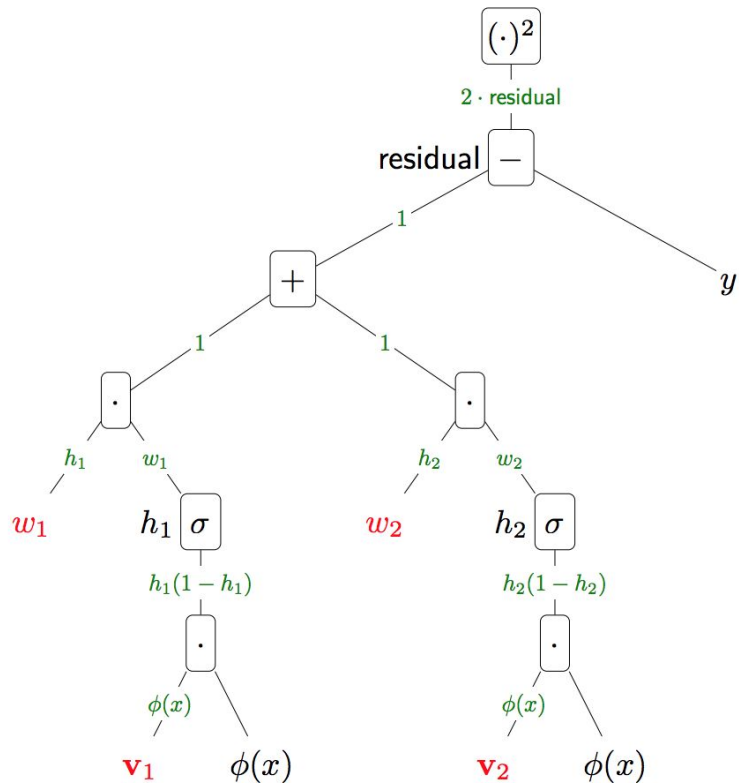
$$-1[\text{margin} < 1]\phi(x)y$$

Neural Network

- 2 layer Neural Network



$$\text{Loss}(x, y, \mathbf{w}) = \left(\sum_{j=1}^k w_j \sigma(\mathbf{v}_j \cdot \phi(x)) - y \right)^2$$



Other Types of Optimizers

Apart from gradient descent there are other optimizers widely used-

- **Adagrad** - Adagrad adapts the learning rate specifically to individual features: that means that some of the weights in your dataset will have different learning rates than others.
- **RMSProp** - RMSProp is Root Mean Square Propagation. It was devised by Geoffrey Hinton. RMSProp tries to resolve Adagrad's radically diminishing learning rates by using a moving average of the squared gradient.
- **Adam** - Adam stands for adaptive moment estimation, and is another way of using past gradients to calculate current gradients. A combination of RMSProp and Adagrad. Widely used in Computer Vision tasks.

Training a Neural Network

1. Pick a network architecture

of input units = # of features

of output units = # of classes

2. Randomly initialize weights
3. Implement forward propagation to get $h(x)$ for any x
4. Compute cost function
5. Use gradient descent/optimizer with backprop to fit the network

Resources-

<https://www.youtube.com/watch?v=aircAruvnKk>

<https://www.youtube.com/watch?v=llg3gGewQ5U>

<https://towardsdatascience.com/activation-functions-and-its-types-which-is-better-a9a5310cc8f>

<http://neuralnetworksanddeeplearning.com/>

<https://playground.tensorflow.org>

Stanford CS229 course