

# Data X

Trees, Entropy, Forests, Boosting  
Data, Signals, and Systems

Ikhlaz Sidhu  
Chief Scientist & Founding Director,  
Sutardja Center for Entrepreneurship & Technology  
IEOR Emerging Area Professor Award, UC Berkeley

# Decision Tree Illustration

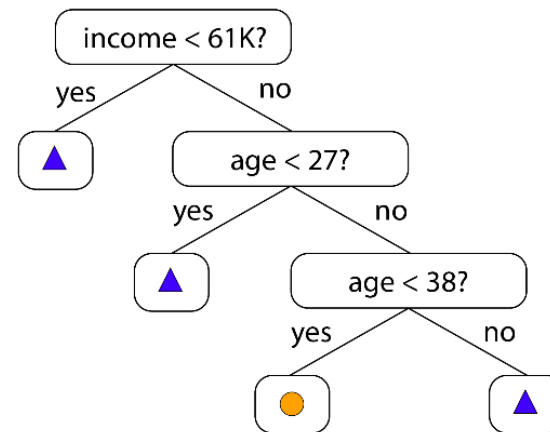
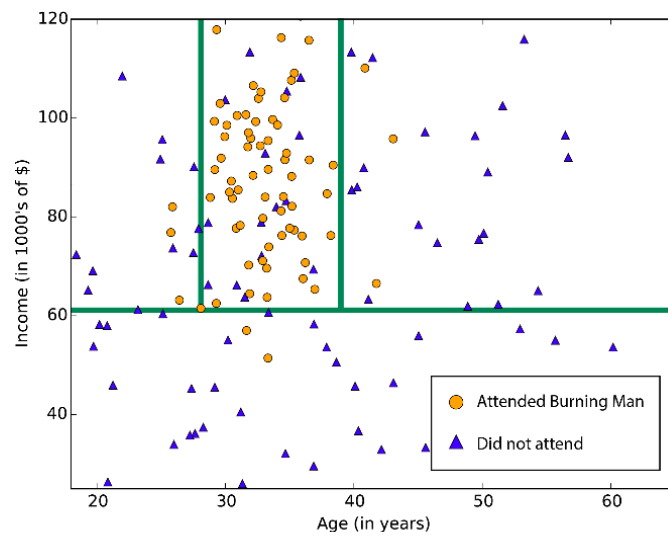
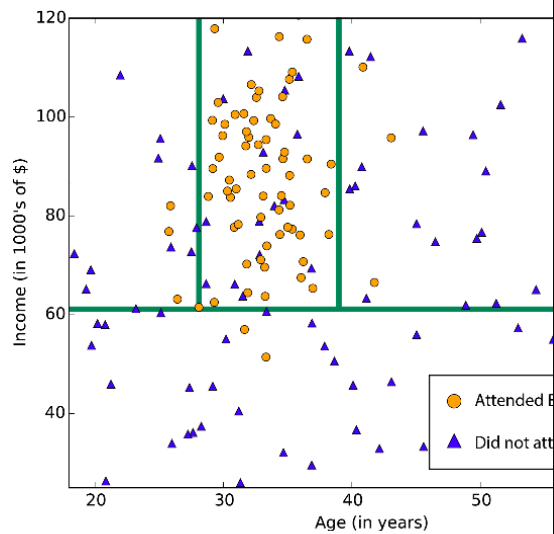


Illustration Source: <https://docs.microsoft.com/en-us/azure/machine-learning/machine-learning-algorithm-choice>

DataX

## Decision Tree Illustration



```
from sklearn import tree
```

```
decision_tree = DecisionTreeClassifier()  
decision_tree.fit(X_train, Y_train)  
Y_pred = decision_tree.predict(X_test)
```

```
# Error
```

```
acc_decision_tree =  
round(decision_tree.score(X_train, Y_train) *  
100, 2)  
acc_decision_tree
```

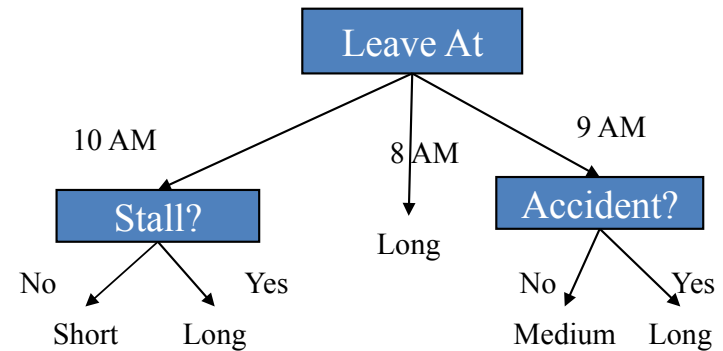
```
# or compare Y_pred with Y_test
```

Illustration Source: <https://docs.microsoft.com/en-us/machine-learning/sklearn/decision-trees>

Data<sup>X</sup>

## How Are Trees Created?

	Attributes				Target
	Hour	Weather	Accident	Stall	Commute
D1	8 AM	Sunny	No	No	Long
D2	8 AM	Cloudy	No	Yes	Long
D3	10 AM	Sunny	No	No	Short
D4	9 AM	Rainy	Yes	No	Long
D5	9 AM	Sunny	Yes	Yes	Long
D6	10 AM	Sunny	No	No	Short
D7	10 AM	Cloudy	No	No	Short
D8	9 AM	Rainy	No	No	Medium
D9	9 AM	Sunny	Yes	No	Long
D10	10 AM	Cloudy	Yes	Yes	Long
D11	10 AM	Rainy	No	No	Short
D12	8 AM	Cloudy	Yes	No	Long
D13	9 AM	Sunny	No	No	Medium



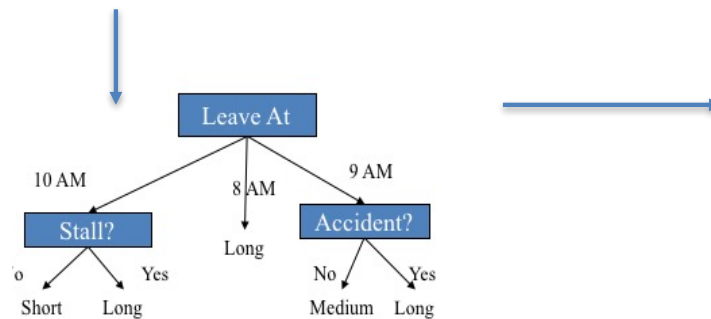
The choice of which feature/attribute to split depends on "entropy"

Example from Jeff Story, source Internet

Data<sup>x</sup>

# How Are Trees Created?

	Attributes				Target
	Hour	Weather	Accident	Stall	Commute
D1	8 AM	Sunny	No	No	Long
D2	8 AM	Cloudy	No	Yes	Long
D3	10 AM	Sunny	No	No	Short
D4	9 AM	Rainy	Yes	No	Long
D5	9 AM	Sunny	Yes	Yes	Long
D6	10 AM	Sunny	No	No	Short
D7	10 AM	Cloudy	No	No	Short
D8	9 AM	Rainy	No	No	Medium
D9	9 AM	Sunny	Yes	No	Long
D10	10 AM	Cloudy	Yes	Yes	Long
D11	10 AM	Rainy	No	No	Short
D12	8 AM	Cloudy	Yes	No	Long
D13	9 AM	Sunny	No	No	Medium



Prediction can be coded

All features might not even be used (Occam's Razor)

```

if hour == 8am
    commute time = long
else if hour == 9am
    if accident == yes
        commute time = long
    else
        commute time = medium
else if hour == 10am
    if stall == yes
        commute time = long
    else
        commute time = short
  
```

Example from Jeff Story, source Internet

DataX

## A Table

Y	X1	X2	X3
0	0	0	1
0	1	0	1
0	0	0	1
0	1	1	1
1	0	1	0
1	1	1	0
1	0	1	0
1	1	0	0

$P(Y=1)$   
= .5

$P(X1=1)$   
= .5

$P(X2=1)$   
= .5

$P(X3=1)$   
= .5

Y = Result output

If you want an efficient tree, which  
Variable (X1, X2, or X3) will you use as the first  
branch of the tree.

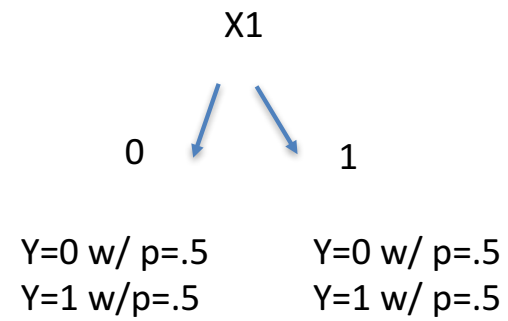
Data<sup>X</sup>



## A Table

Y	X1	X2	X3
0	0	0	1
0	0	0	1
1	0	1	0
1	0	1	0
0	1	0	1
0	1	1	1
1	1	1	0
1	1	0	0
Sorted X1 and then Y			

Y=1 w/p=.5

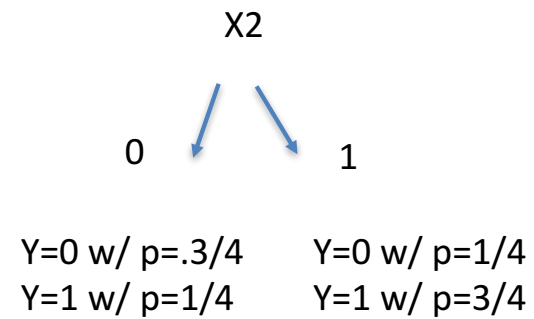


Data<sup>X</sup>

## A Table

	Y	X1	X2	X3
	0	0	0	1
	0	1	0	1
	0	0	0	1
	1	1	0	0
	0	1	1	1
	1	0	1	0
	1	1	1	0
	1	0	1	0
sorted X2 and then by Y				

Y=1 w/p=.5

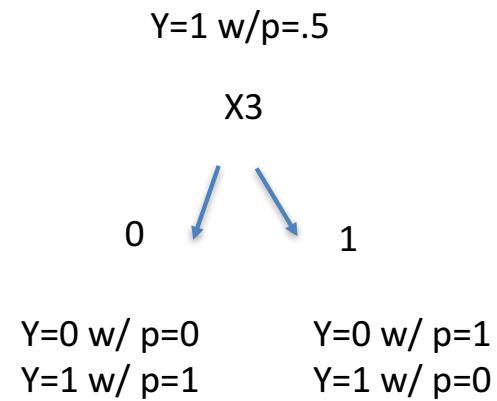


Data<sup>X</sup>



## A Table

Y	X1	X2	X3
1	0	1	0
1	1	1	0
1	0	1	0
1	1	0	0
0	0	0	1
0	1	0	1
0	0	0	1
0	1	1	1
Sorted X3 and then by Y			



Data<sup>X</sup>

How much information is in a 'data value' from column, eg 'Y'?

$$P(Y=1) = .5$$

Y	X1	X2	X3
0	0	0	1
0	1	0	1
0	0	0	1
0	1	1	1
1	0	1	0
1	1	1	0
1	0	1	0
1	1	0	0

How much information is in an observed value  
In the Y column?

$$Y = \begin{cases} 1 \text{ w } p = 1/2 \\ 0 \text{ with } p = 1/2 \end{cases} \longrightarrow \begin{array}{l} \text{Suppose} \\ \text{outcome} \\ = 1 \end{array}$$

$$Y = \begin{cases} 1 \text{ w } p = 1 \\ 0 \text{ w } p = 0 \end{cases} \longrightarrow \begin{array}{l} \text{Suppose} \\ \text{outcome} \\ = 1 \end{array}$$

Data<sup>x</sup>

## We can use “Entropy” as a Measure of Information

$$H(x) = \sum_x p(x) \log \left( \frac{1}{p(x)} \right)$$

A Measure of a distribution  $P(X)$

Y	X1	X2	X3
0	0	0	1
0	1	0	1
0	0	0	1
0	1	1	1
1	0	1	0
1	1	1	0
1	0	1	0
1	1	0	0

Examples:

$X = 1$  with  $p = .5$

$$H(X) = .5 \log_2 (1/.5) + .5 \log_2 (1/.5) = 1 \text{ bits}$$

$X = 1$  with  $p = 1/4$

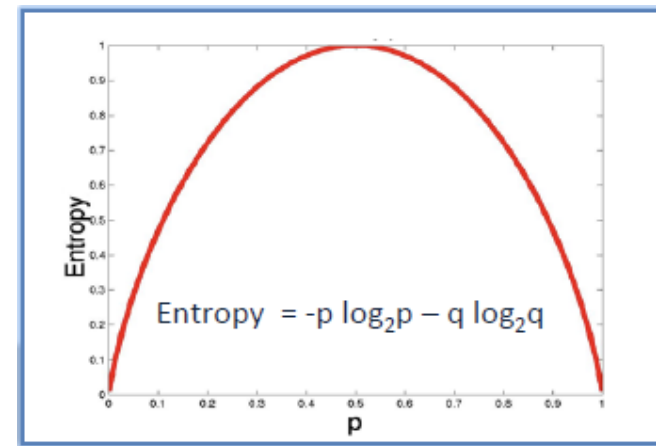
$$H(X) = 1/4 \log_2 (4) + 3/4 \log_2 (4/3) = 0.811 \text{ bits}$$

$X = 1$  with  $p = 1/10$

$$H(X) = .1 \log_2 (10) + 0.9 \log_2 (10/9/3) = .46 \text{ bits}$$

$X=1$  w  $p = 0$

$$H(X) = 0 \log_2 (1/0^+) + 1 \log_2 (1) = 0 \text{ bits}$$

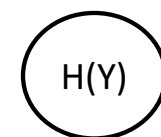


Data<sup>X</sup>

## How Do We Construct a Tree. Example ID3 Algorithm

- Find  $H(Y)$  = Entropy of Outcome
- Then try each possible feature.  
Variable:  $X_1, X_2$ , and  $X_3$
- Calculate the Conditional Entropy of  $Y$  given a choice of either  $X_1, X_2$  or  $X_3$   
 $= H(Y|X)$
- Information Gained =  
•  $I(Y;X) = H(Y) - H(Y|X)$
- Choose the variable  $X_1, X_2, X_3$  that leads to the most information gained.

$Y=1$  w/ $p=.5$



Entropy of Outcome

$X_1 = 0$

$X_1 = 1$

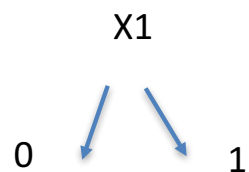
Which Variable to Split?  
 $X_1, X_2$  or  $X_3$ ?

Distribution	$P(Y X_1=0)$	$P(Y X_1=1)$
Conditional Entropy	$H(Y X_1=0)$	$H(Y X_1=1)$

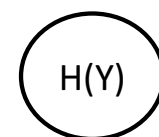
Data<sup>X</sup>

## How Do We Construct a Tree. Example ID3 Algorithm

Y	X1	X2	X3
0	0	0	1
0	0	0	1
1	0	1	0
1	0	1	0
0	1	0	1
0	1	1	1
1	1	1	0
1	1	0	0
Sorted X1 and then Y			



Y=0 w/ p=.5    Y=0 w/ p=.5  
Y=1 w/p=.5    Y=1 w/ p=.5



Entropy of Outcome  
 $H(Y) = 1$

X1 = 0  
W p = 1/2

X1 = 1

Which Variable to Split?  
Let's Try X1

$$P(Y=1|X1=0) = 0.5$$

$$P(Y=1|X1=1) = 0.5$$

$$H(Y|X1=0) = 1$$

$$H(Y|X1=1) = 1$$

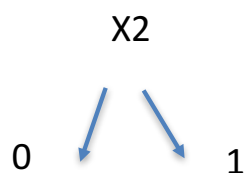
$$H(Y|X) = 1 \times .5 + 1 \times .5 = 1.0$$

$$\text{Info Gained} = H(Y) - H(Y|X) = 1 - 1 = 0$$

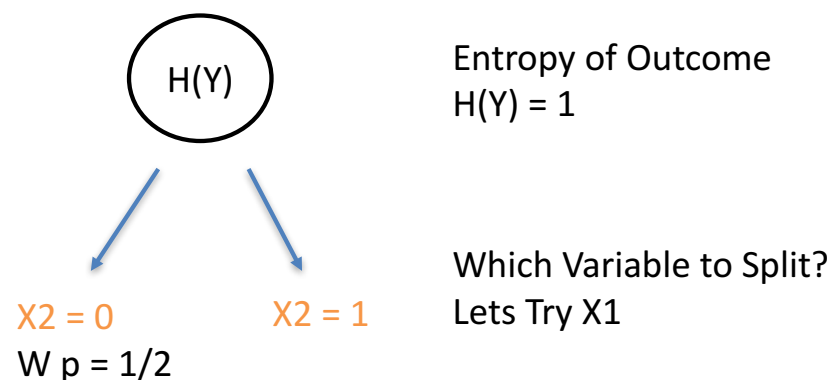
Data X

## How Do We Construct a Tree. Example ID3 Algorithm

Y	X1	X2	X3
0	0	0	1
0	1	0	1
0	0	0	1
1	1	0	0
0	1	1	1
1	0	1	0
1	1	1	0
1	0	1	0
sorted X2 and then by Y			



Y=0 w/  $p=.3/4$       Y=0 w/  $p=1/4$   
 Y=1 w/  $p=1/4$       Y=1 w/  $p=3/4$



Entropy of Outcome  
 $H(Y) = 1$

$$P(Y=1|X2=0) = 1/4$$

$$P(Y=1|X2=1) = 3/4$$

$$H(Y|X1=0) = \frac{1}{4}\log 4 + \frac{3}{4}\log \frac{4}{3} = .81$$

$$H(Y|X1=1) = \frac{3}{4}\log \frac{4}{3} + \frac{1}{4}\log 4 = 0.81$$

$$H(Y|X) = 0.81 \times 0.5 + 0.81 \times 0.5 = 0.81$$

$$\text{Info Gained} = H(Y) - H(Y|X) = 1 - 0.81 = 0.19$$

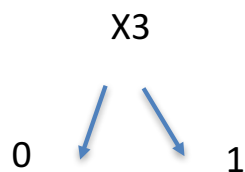
Data X



## How Do We Construct a Tree. Example ID3 Algorithm

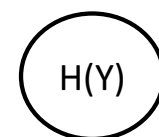
Y	X1	X2	X3
0	0	0	1
0	1	0	1
0	0	0	1
0	1	1	1
1	0	1	0
1	1	1	0
1	0	1	0
1	1	0	0

Y=1 w/p=.5



Y=1 w/ p=1

Y=1 w/ p=0



Entropy of Outcome  
 $H(Y) = 1$



$X3 = 0$   
W p = 1/2

$X3 = 1$

Which Variable to Split?  
Let's Try X1

$P(Y=1 | X3=0)$   
 $= 1$

$P(Y=1 | X3=1)$   
 $= 0$

$H(Y | X1=0)$   
 $= 1 \log(1/1) +$   
 $0 \log(1/0)$   
 $= 0$

$H(Y | X1=1)$   
 $= 1 \log(1/1) + 0 \log 1/0$   
 $= 0$

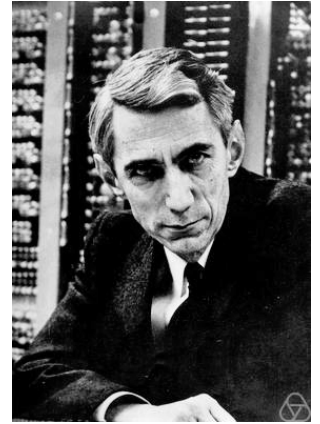
$H(Y | X) = 0 \times 0.5 + 0 \times 0.5 = 0$

Info Gained =  $H(Y) - H(Y | X) = 1 - 0 = 1$

Data X

## Lets Understand Entropy

Claude Shannon,  
"the father of  
[information theory](#)"



Entropy is a measure of Information in a source  
(or transmitted, lost or received)

Example:

- 4 Symbols: A, B, C, D
- $P(A) = 6/12$ ,  $P(B) = 3/12$ ,  $P(C) = 2/12$ ,  $P(D) = 1/12$

Possibly transmitted in this sequence: AABA ABBC CAAD ... → “?”

Q: How many bits of information is in A, B, C, D and the “next missing random symbol”



## Lets First Understand Entropy

Entropy is a measure of Information in a source (or transmitted, lost or received)

Example:

- 4 Symbols: A, B, C, D
- $P(A) = 6/12$ ,  $P(B) = 3/12$ ,  $P(C) = 2/12$ ,  $P(D) = 1/12$

Possibly transmitted in this sequence: AABA ABBC CAAD ...  $\rightarrow$  “?”

Q: How many bits of information is in A, B, C, D and the “next missing random symbol”

**Answer relates to Entropy**

- Info Content:  $h(p) = p \log_2(1/p)$  for a symbol with probability  $p$
- $h(A) = \log_2(2) = 1$  bit,  $h(B) = \log_2(4) = 2$  bits,  $h(C) = \log_2(6)$ ,  $h(D) = \log_2(12)$
- $H(S) = \frac{1}{2} \log_2(2) + \frac{1}{4} \log_2(4) + \frac{1}{6} \log_2(6) + \frac{1}{12} \log_2(12) = 1.72$  bits

Information:

- Increases with less likelihood
- Additive for combinations  
Options A x B  $\rightarrow H(A) + H(B)$
- An event with  $p = 0$  has no info

$$H(x) = \sum_x p(x) \log \left( \frac{1}{p(x)} \right)$$



## Entropy Related to the Information Contained in the Symbol (or Table)

Entropy is a measure of Information in a source (or transmitted, lost or received)

Example:

- $P(A) = 6/12$
- $P(B) = 3/12$
- $P(C) = 2/12$
- $P(D) = 1/12$
- Possibly transmitted in this sequence: AABA ABBC CAAD ...  $\rightarrow$  “?”

Source Coding Theorem:

The  $H(S)$  is the smallest codeword length that is theoretically possible for signal ‘S’

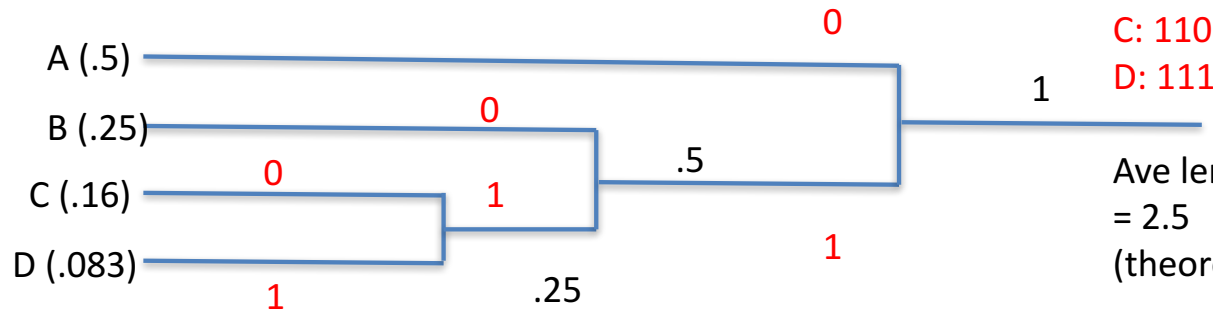
Theoretically smallest code is 1.72 bits per symbol



## Sample Coding Schemes

Example:

- 4 Symbols: A, B, C, D
- Suppose each is “encoded” in 4 bits, ie A = ‘0001’
- $P(A) = 6/12$ ,  $P(B) = 3/12$ ,  $P(C) = 2/12$ ,  $P(D) = 1/12$
- We could code in 2 bits (00, 01, 10, 11)
- Most famous prefix free coding is called a Huffman Code:



A prefix free code:

A: 0

B: 10

C: 110

D: 111

Ave length:  $0.5(1) + 0.25(2) + 0.16(3) + 0.083(3)$   
 $= 2.5$

(theoretically best would be 1.72)

DataX

## Examples with Tables, Trees and Entropy



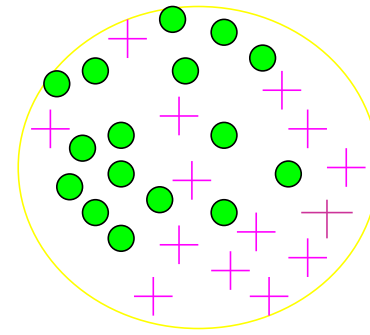


## Entropy Example

- Entropy = 
$$\sum_i -p_i \log_2 p_i$$

$p_i$  is the probability of class  $i$

Compute it as the proportion of class  $i$  in the set.



16/30 are green circles; 14/30 are pink crosses

$\log_2(16/30) = -.9$ ;  $\log_2(14/30) = -1.1$

Entropy =  $-(16/30)(-.9) - (14/30)(-1.1) = .99$

Source Internet

DataX

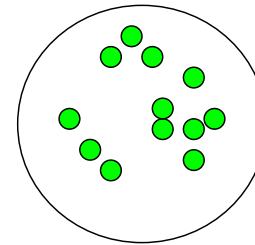
## 2-Class Cases:

- What is the entropy of a group in which all examples belong to the same class?

–  $\text{entropy} = -1 \log_2 1 = 0$

not a good training set for learning

**Minimum  
impurity**

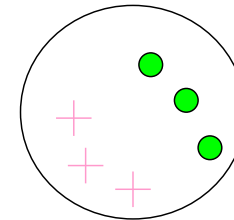


- What is the entropy of a group with 50% in either class?

–  $\text{entropy} = -0.5 \log_2 0.5 - 0.5 \log_2 0.5 = 1$

good training set for learning

**Maximum  
impurity**

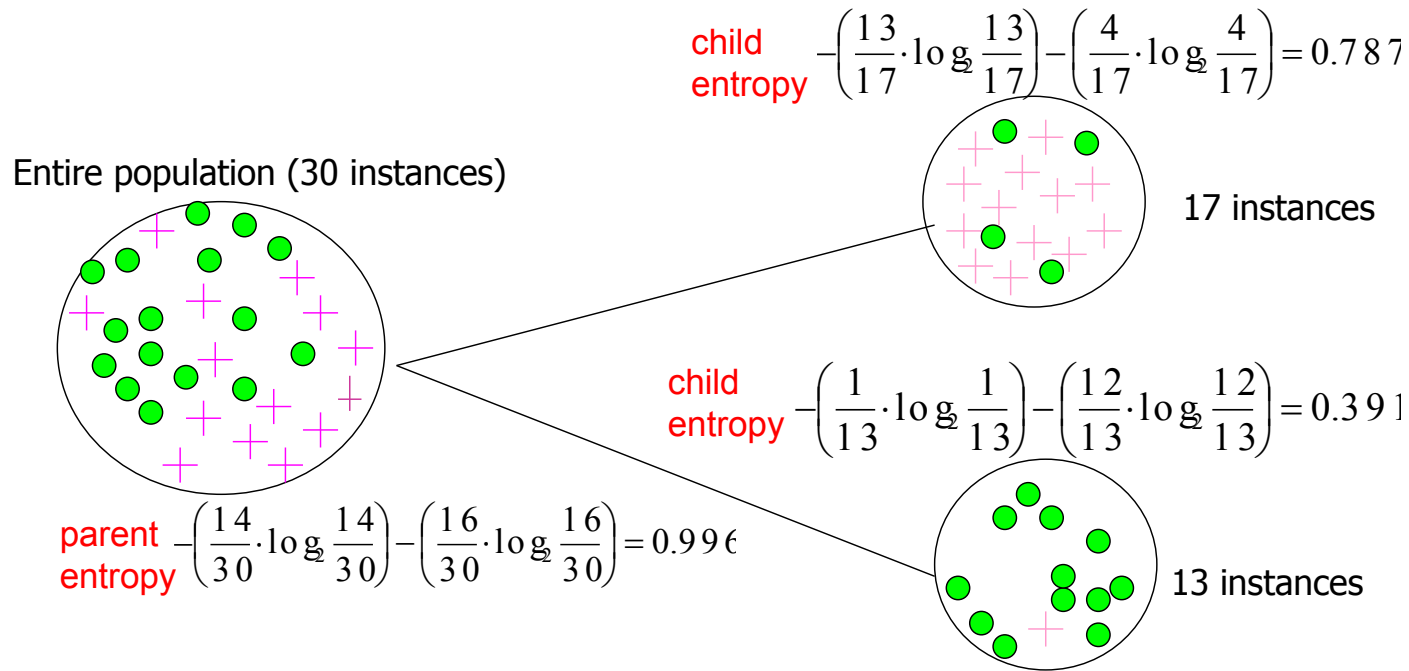


Source Internet

Data<sup>x</sup>

## Calculating Information Gain

**Information Gain** = entropy(parent) – [average entropy(children)]



**(Weighted) Average Entropy of Children** =  $\left(\frac{17}{30} \cdot 0.787\right) + \left(\frac{13}{30} \cdot 0.391\right) = 0.615$

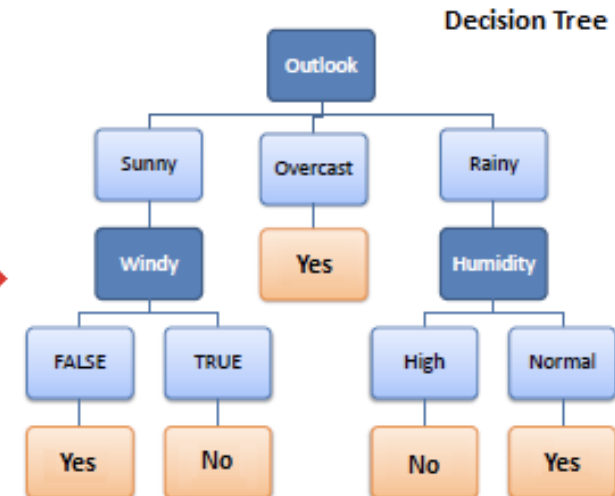
**Information Gain** = 0.996 - 0.615 = 0.38 for this split

7

Data<sup>x</sup>

## Golf Example

Predictors				Target
Outlook	Temp.	Humidity	Windy	Play Golf
Rainy	Hot	High	False	No
Rainy	Hot	High	True	No
Overcast	Hot	High	False	Yes
Sunny	Mild	High	False	Yes
Sunny	Cool	Normal	False	Yes
Sunny	Cool	Normal	True	No
Overcast	Cool	Normal	True	Yes
Rainy	Mild	High	False	No
Rainy	Cool	Normal	False	Yes
Sunny	Mild	Normal	False	Yes
Rainy	Mild	Normal	True	Yes
Overcast	Mild	High	True	Yes
Overcast	Hot	Normal	False	Yes
Sunny	Mild	High	True	No



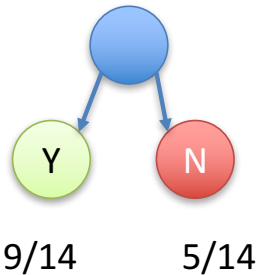
Start with table of experiences

- Distribution:  $P(\text{golf}) = 9 \text{ yes}, 5 \text{ no}$  14 observations
- We calculate entropy of distribution: playing golf
- Need to calculate entropy of each next possible decision (Outlook, Temp, Humidity, and Windy)
- Information gained is  $E(\text{Golf}) - E(\text{Subset}, \text{Golf})$ ,
- We choose the next branch to be the decision with the highest information gained for the next branch

DataX

# Golf Example

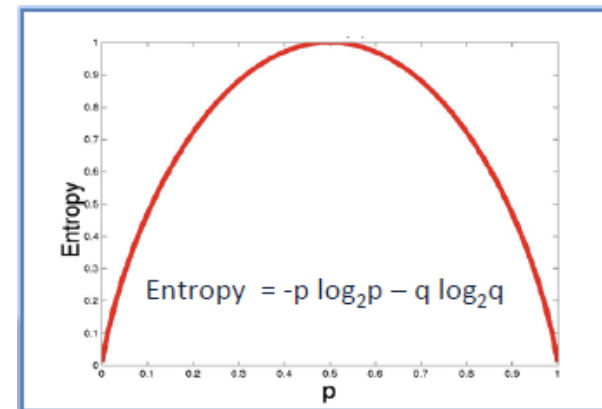
Top Level  
Play Golf  
Distribution



$$E(G) = 9/14 \log_2 1/(9/14) + 5/14 \log_2 1/(5/14) = 0.94$$

- Distribution:  $P(\text{golf}) = 9 \text{ yes}, 5 \text{ no} / 14 \text{ observations}$
- We calculate entropy of distribution: playing golf
- Need to calculate entropy of each next possible decision (Outlook, Temp, Humidity, and Windy)
- Information gained is  $E(\text{Golf}) - E(\text{Subset}, \text{Golf})$ ,
- We choose the next branch to be the decision with the highest information gained for the next branch

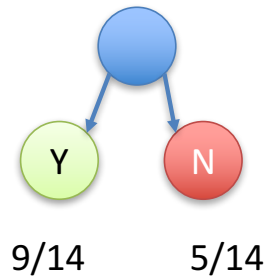
Predictors				Target
Outlook	Temp	Humidity	Windy	Play Golf
Rainy	Hot	High	False	No
Rainy	Hot	High	True	No
Overcast	Hot	High	False	Yes
Sunny	Mild	High	False	Yes
Sunny	Cool	Normal	False	Yes
Sunny	Cool	Normal	True	No
Overcast	Cool	Normal	True	Yes
Rainy	Mild	High	False	No
Rainy	Cool	Normal	False	Yes
Sunny	Mild	Normal	False	Yes
Rainy	Mild	Normal	True	Yes
Overcast	Mild	High	True	Yes
Overcast	Hot	Normal	False	Yes
Sunny	Mild	High	True	No



$$\text{Entropy} = -0.5 \log_2 0.5 - 0.5 \log_2 0.5 = 1$$

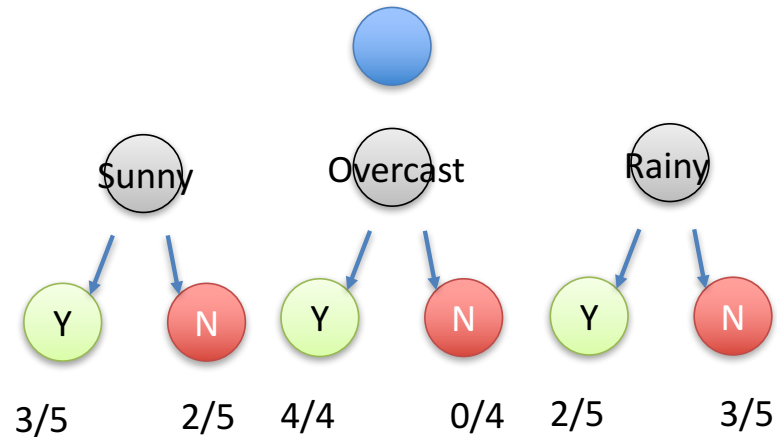
Data<sup>x</sup>

Top Level  
Play Golf  
Distribution



$$\begin{aligned}
 H(G) &= 9/14 \log_2 1/(9/14) \\
 &\quad + 5/14 \log_2 1/(5/14) \\
 &= 0.94
 \end{aligned}$$

Next Level: Play Golf with Outlook



$$\begin{aligned}
 H(G,OL) &= \\
 &P(\text{Sunny}) H(G,OL) + P(\text{OV})H(G,OV) + P(\text{Rainy}) H(G,\text{Rainy}) \\
 &= 5/14 * 0.971 + 4/14 * 0 + 5/14 * .971 = 0.693
 \end{aligned}$$

$$\text{Info Gained} = H(G) - H(G,OL) = 0.94 - 0.693 = 0.247$$

Do the same with **Temp (.29)**, Windy(.048), and Humidity(.152):  
Choose next node to be the one with the most info gained

Data<sup>x</sup>



# Random Forrest



## Trees Can be Extended with Bagging

Explain  
bagging and  
Random  
Forrest

```
from sklearn.ensemble import RandomForestClassifier

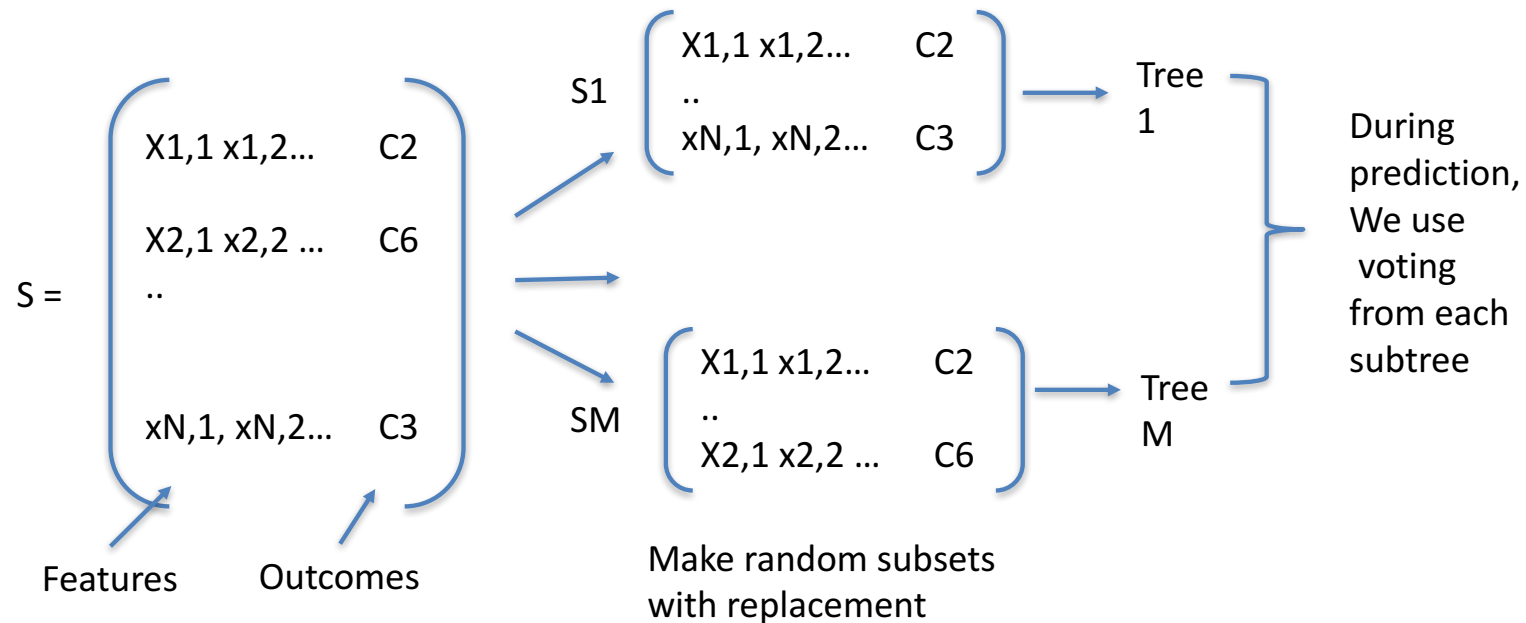
random_forest =
RandomForestClassifier(n_estimators=1000)
random_forest.fit(X_train, Y_train)
Y_pred = random_forest.predict(X_test)
random_forest.score(X_train, Y_train)

# Error
acc_random_forest = round(random_forest.score(X_train,
Y_train) * 100, 2)
acc_random_forest

# or compare Y_pred with Y_test
```

Data<sup>X</sup>

## Random Forest – A type of bagging/ensemble approach



Advantages: One of most accurate  
Efficient prediction over large data

Disadvantages: Overfit and Training time

Data<sup>x</sup>

## Our experiment with the Titanic Data Set

	<b>Model</b>	<b>Score</b>
	Random Forest	86.76
	Decision Tree	86.76
	KNN	84.74
	Support Vector Machines	83.84
	Logistic Regression	80.36
	Linear SVC	79.01
	Perceptron	78.00
	Naive Bayes	72.28
	Stochastic Gradient Decent	72.28



More Accuracy  
Generally more training time  
More risk of overfitting

Less Accuracy  
Generally less computation

Data<sup>x</sup>

# Boosting

Data<sup>X</sup>

A decorative horizontal bar at the bottom of the slide. It features a dark background with a pattern of glowing blue and white binary digits (0s and 1s) arranged in a way that suggests data flow or a digital landscape. The word "DataX" is prominently displayed in a white, serif font on the left side of the bar.

## Boosting as in AdaBoost

- Motivation - a procedure that combines the outputs of many “weak” classifiers to produce a powerful “committee”
- A machine learning algorithm
- Perform supervised learning
- Increments improvement of learned function
- Forces the weak learner to generate new hypotheses that make less mistakes on “harder” parts.

- Freund & Schapire (1995) – **AdaBoost**

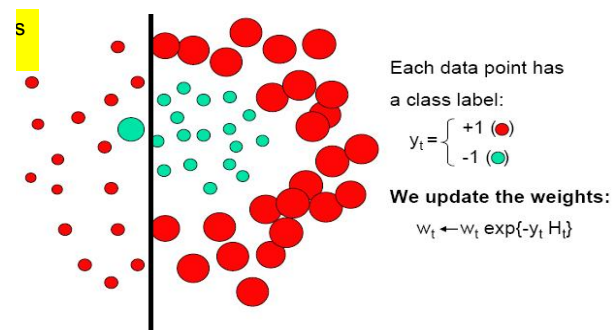
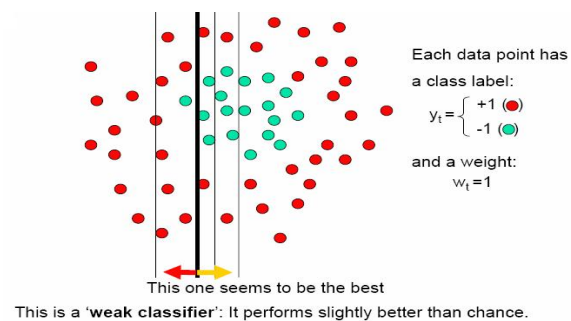
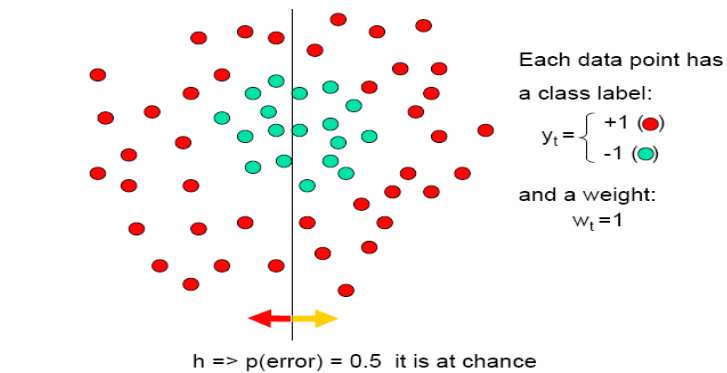
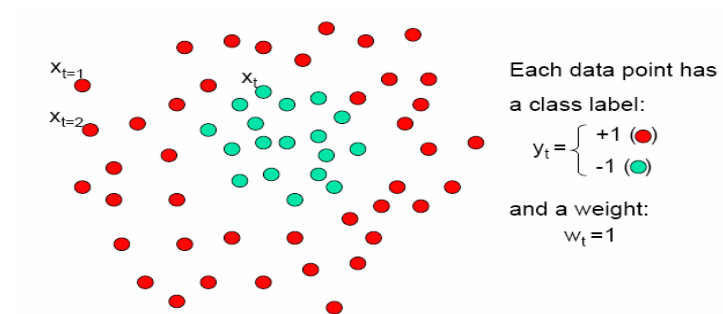
- strong practical advantages over previous boosting algorithms

Benk Erika

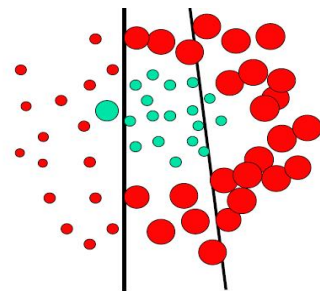
Kelemen Zsolt

Data<sup>x</sup>





Data<sup>X</sup>



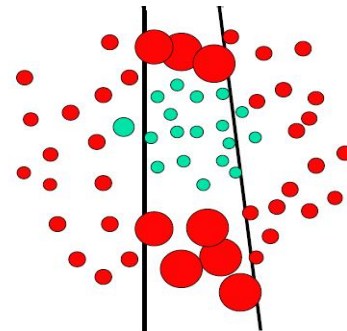
Each data point has  
a class label:

$$y_i = \begin{cases} +1 & (\text{red}) \\ -1 & (\text{green}) \end{cases}$$

We update the weights:

$$w_i \leftarrow w_i \exp\{-y_i H_i\}$$

We set a new problem for which the previous weak classifier performs at chance again

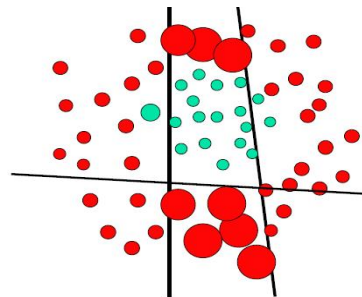


Each data point has  
a class label:

$$y_i = \begin{cases} +1 & (\text{red}) \\ -1 & (\text{green}) \end{cases}$$

We update the weights:

$$w_i \leftarrow w_i \exp\{-y_i H_i\}$$

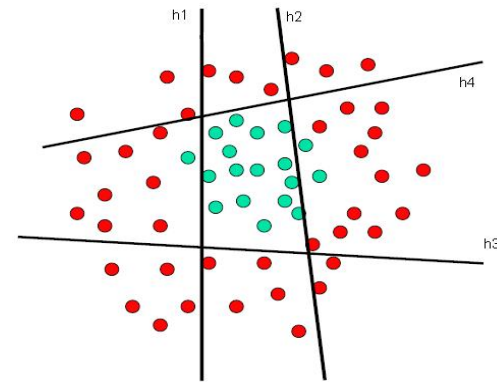


Each data point has  
a class label:

$$y_i = \begin{cases} +1 & (\text{red}) \\ -1 & (\text{green}) \end{cases}$$

We update the weights:

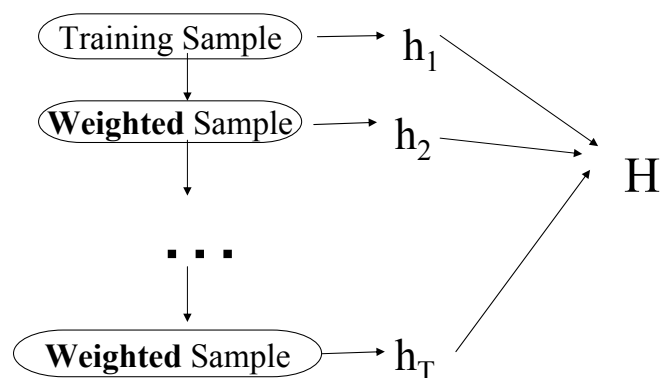
$$w_i \leftarrow w_i \exp\{-y_i H_i\}$$



A strong (or linear) classifier is built as the comb

Data<sup>X</sup>

## General Concept of Boosting



- Train a set of weak hypotheses:  $h_1, \dots, h_T$ .
- The combined hypothesis  $H$  is a **weighted** majority vote of the  $T$  weak hypotheses.
  - Each hypothesis  $h_t$  has a weight  $\alpha_t$ .

$$H(x) = \text{sign}\left(\sum_{t=1}^T \alpha_t h_t(x)\right)$$

Benk Erika  
Kelemen Zsolt

Data<sup>x</sup>

## Boosting as in AdaBoost

# Boosting

- Binary classification problem
- Training data:

$(x_1, y_1), \dots, (x_m, y_m)$ , where  $x_i \in X, y_i \in Y = \{-1, 1\}$

- $D_t(i)$ : the weight of  $x_i$  at round  $t$ .  $D_1(i) = 1/m$ .
- A learner  $L$  that finds a weak hypothesis  $h_t: X \rightarrow Y$  given the training set and  $D_t$
- The error of a weak hypothesis  $h_t$ :

$$\varepsilon_t = \Pr_{i \sim D_t} [h_t(x_i) \neq y_i] = \sum_{i: h_t(x_i) \neq y_i} D_t(i)$$

Benk Erika  
Kelemen Zsolt

Data<sup>X</sup>

## AdaBoost Algorithm

- **Given**  $(x_1, y_1), \dots, (x_m, y_m)$  where  $x_i \in X$ ,  $y_i \in \{-1, +1\}$
- **Initialise** weights  $D_1(i) = 1/m$
- **Iterate**  $t=1, \dots, T$ :
  - Train weak learner using distribution  $D_t$
  - Get weak classifier:  $h_t: X \rightarrow \mathbb{R}$
  - Choose  $\alpha_t \in \mathbb{R}$
  - Update:  $D_{t+1}(i) = \frac{D_t(i) \exp(-\alpha_t y_i h_t(x_i))}{Z_t}$
- where  $Z_t$  is a normalization factor (chosen so that  $D_{t+1}$  will be a distribution), and  $\alpha_t$ :

$$\alpha_t = \frac{1}{2} \ln \left( \frac{1 - \varepsilon_t}{\varepsilon_t} \right) > 0$$

- **Output** – the final classifier

$$H(x) = \text{sign} \left( \sum_{t=1}^T \alpha_t h_t(x) \right)$$

Benk Erika  
Kelemen Zsolt

Data<sup>X</sup>

## AdaBoost Algorithm

### Discrete AdaBoost - Algorithm

- **Given**  $(x_1, y_1), \dots, (x_m, y_m)$  where  $x_i \in X, y_i \in \{-1, +1\}$
- **Initialise** weights  $D_1(i) = 1/m$
- **Iterate**  $t=1, \dots, T$ :
  - Find  $h_t = \arg \min_{h_j} \epsilon_j$  where  $\epsilon_j = \sum_{i=1}^m D_t(i) \mathbb{I}[h_j(x_i) \neq y_i]$
  - Set

$$\alpha_t = \frac{1}{2} \ln \left( \frac{1 - \epsilon_t}{\epsilon_t} \right)$$

- Update:  $D_{t+1}(i) = \frac{D_t(i) \exp(-\alpha_t y_i h_t(x_i))}{Z_t}$

- **Output** – the final classifier

$$H(x) = \text{sign} \left( \sum_{t=1}^T \alpha_t h_t(x) \right)$$

Benk Erika  
Kelemen Zsolt

Data<sup>X</sup>

# AdaBoost – Pros and Contrasts

## ■ Pros:

- Very simple to implement
- Fairly good generalization
- The prior error need not be known ahead of time

## ■ Contrasts:

- Suboptimal solution
- Can over fit in presence of noise



End of Section

Data<sup>x</sup>